

Hands on **GraphRAG Workshop**

Rabobank Graph Day 2025

Erik Bijl | Neo4j

Martijn Vlak | Rabobank

Who are we?



Erik Bijl

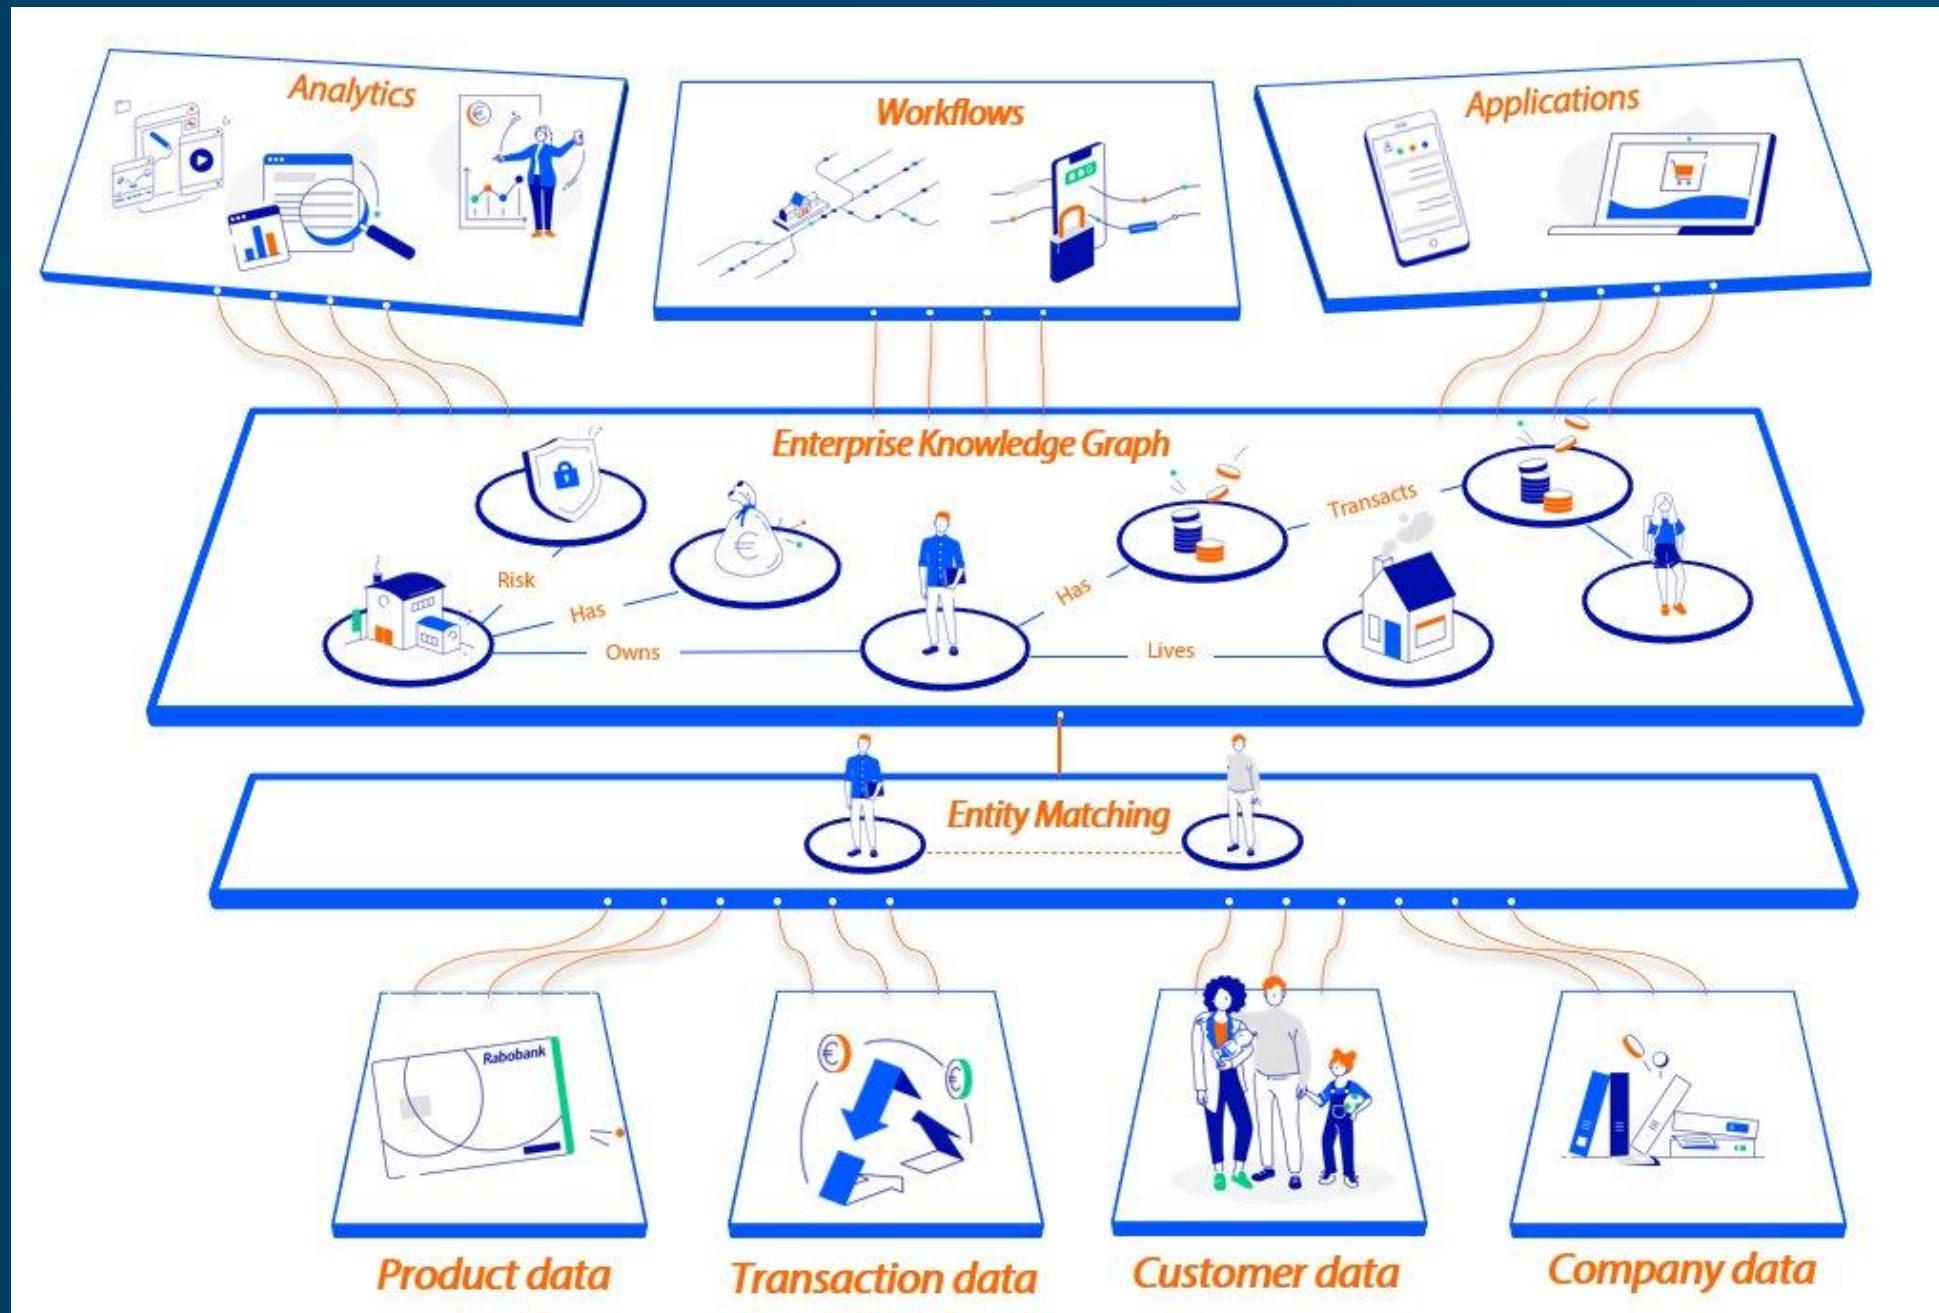
(Erik) - [:LIVES_IN] → (`Zwolle`)
(Erik) - [:WORKS_FOR] → (`Neo4j`)
(Erik) - [:HAS_ROLE] → (`Solution Engineer`)
(Erik) - [:IS_PART_OF] → (`EMEA Field Team`)

Martijn Vlak

(Martijn) - [:LIVES_IN] → (`Leiden`)
(Martijn) - [:WORKS_FOR] → (`Rabobank`)
(Martijn) - [:HAS_ROLE] → (`Data Scientist`)
(Martijn) - [:IS_PART_OF] → (`LinQ-Team`)



Meet LinQ



Agenda

- 1 Workshop objectives**
Setting the stage
- 2 Groups and setup**
Let's get set up
- 3 Modules**
Run the Notebooks
- 4 Wrap up**
Resources and Q&A

Workshop Rules

- Ask questions straight away, this is an interactive session
- Raise your hand if you are stuck
- Slides & notebooks will be shared
- Have fun!

Connect to the notebooks:

- <https://rabobank-jp-one.graphdatabase.ninja>
- Attendeexxx (If you got number 151 => attendee151)



Neo4j Browser: <https://browser.neo4j.io/preview/>

Quick Poll (by show of hands)



Modules

1 Explore the Graph
Run some queries

2 Vector Index
Set up the Vector Search

3 Graph Analytics
Run Graph Algorithms

4 GraphRAG Chatbot
Run a Chatbot on the Graph

5 GraphRAG Agent
Create Agents with Tools

6 Wrap up
Resources and Q&A

Module 1

Explore the Graph: Let's run some queries

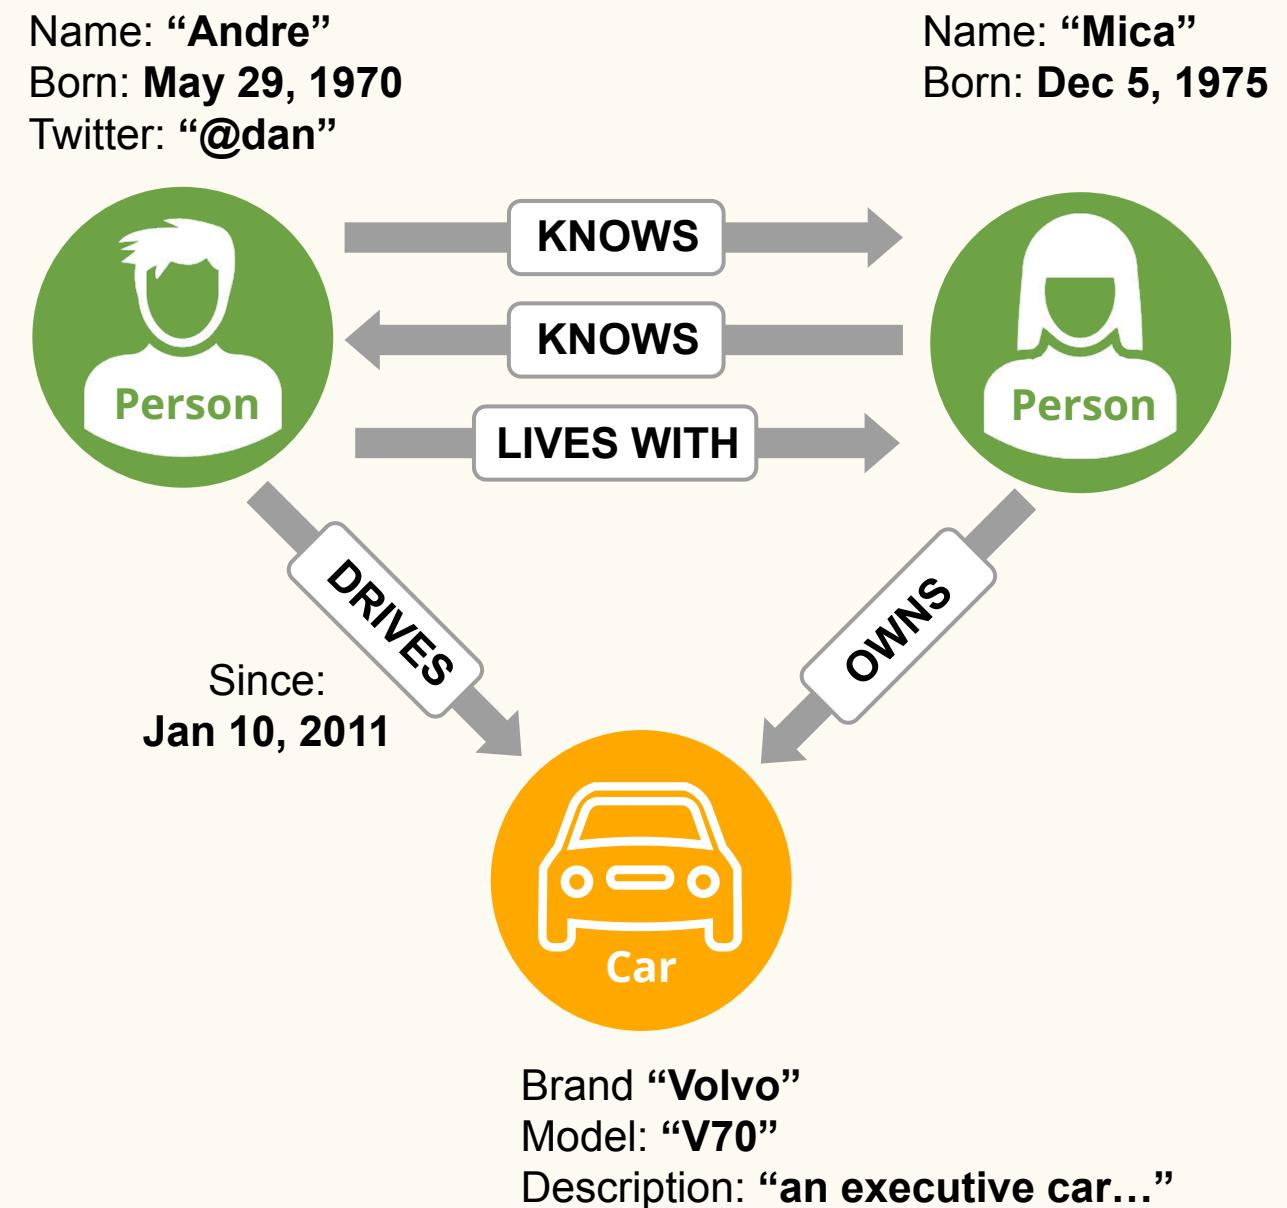
Knowledge Graph = design patterns to organize & access interrelated data

Property Graph Data Model

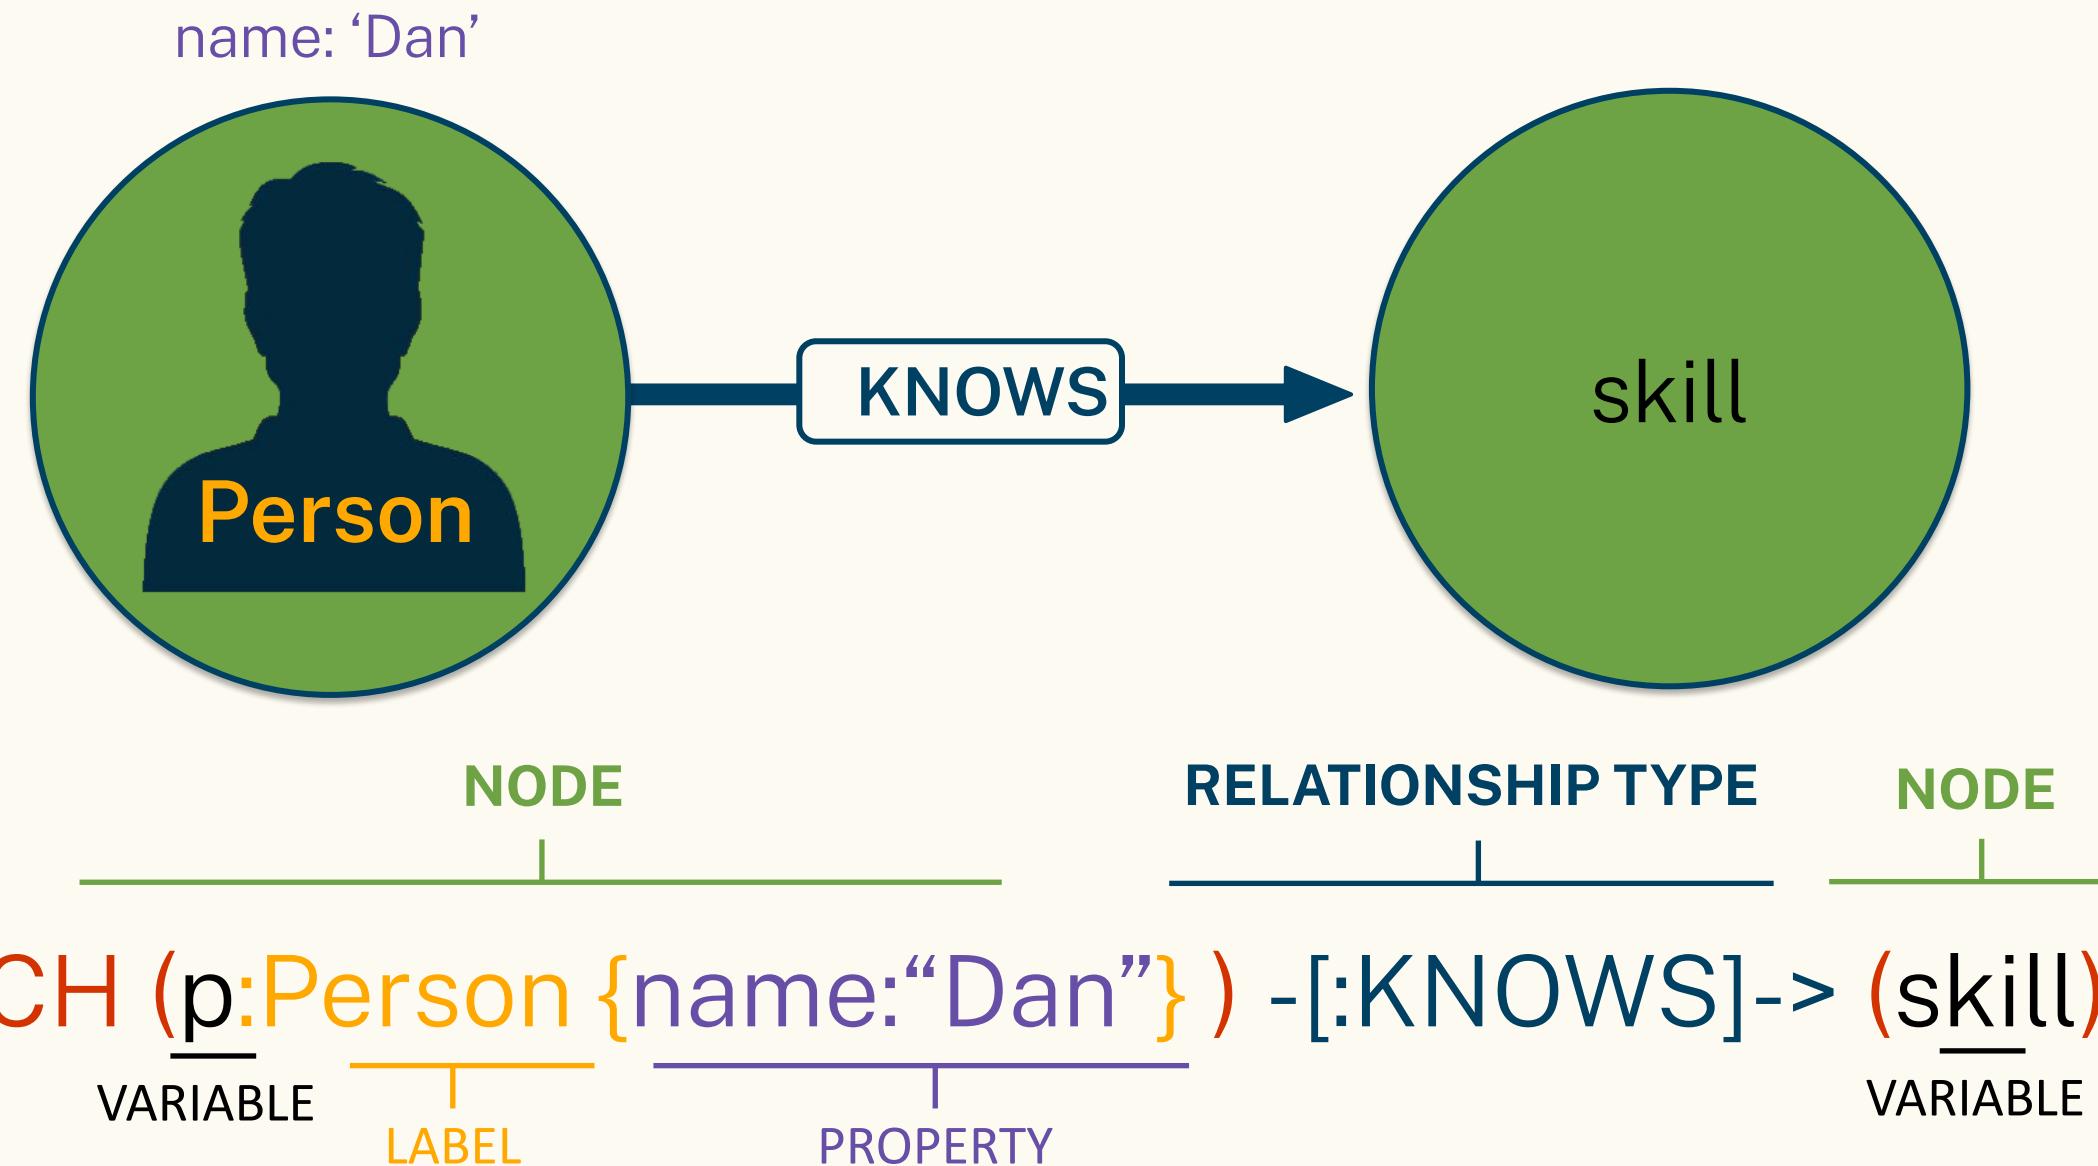
Nodes represent entities in the graph

Relationships represent associations or interactions between nodes

Properties represent attributes of nodes or relationships



Cypher: A Powerful & Expressive Query Language



RETURN p.name as person, skill

From Unstructured to Structured: An example...

RFP Generation GenAI App

Anatomy of an RFP Document

AWS RFP

Intro

About the Company

Content

Financial Result

Content

Objectives

Content

Proposal

Subsection 1

Subsection 1.1

Content

Subsection 2

Content

Anatomy of a Document

AWS RFP

Intro

About the Company

Content

Financial Result

Content

Objectives

Content

Proposal

Subsection 1

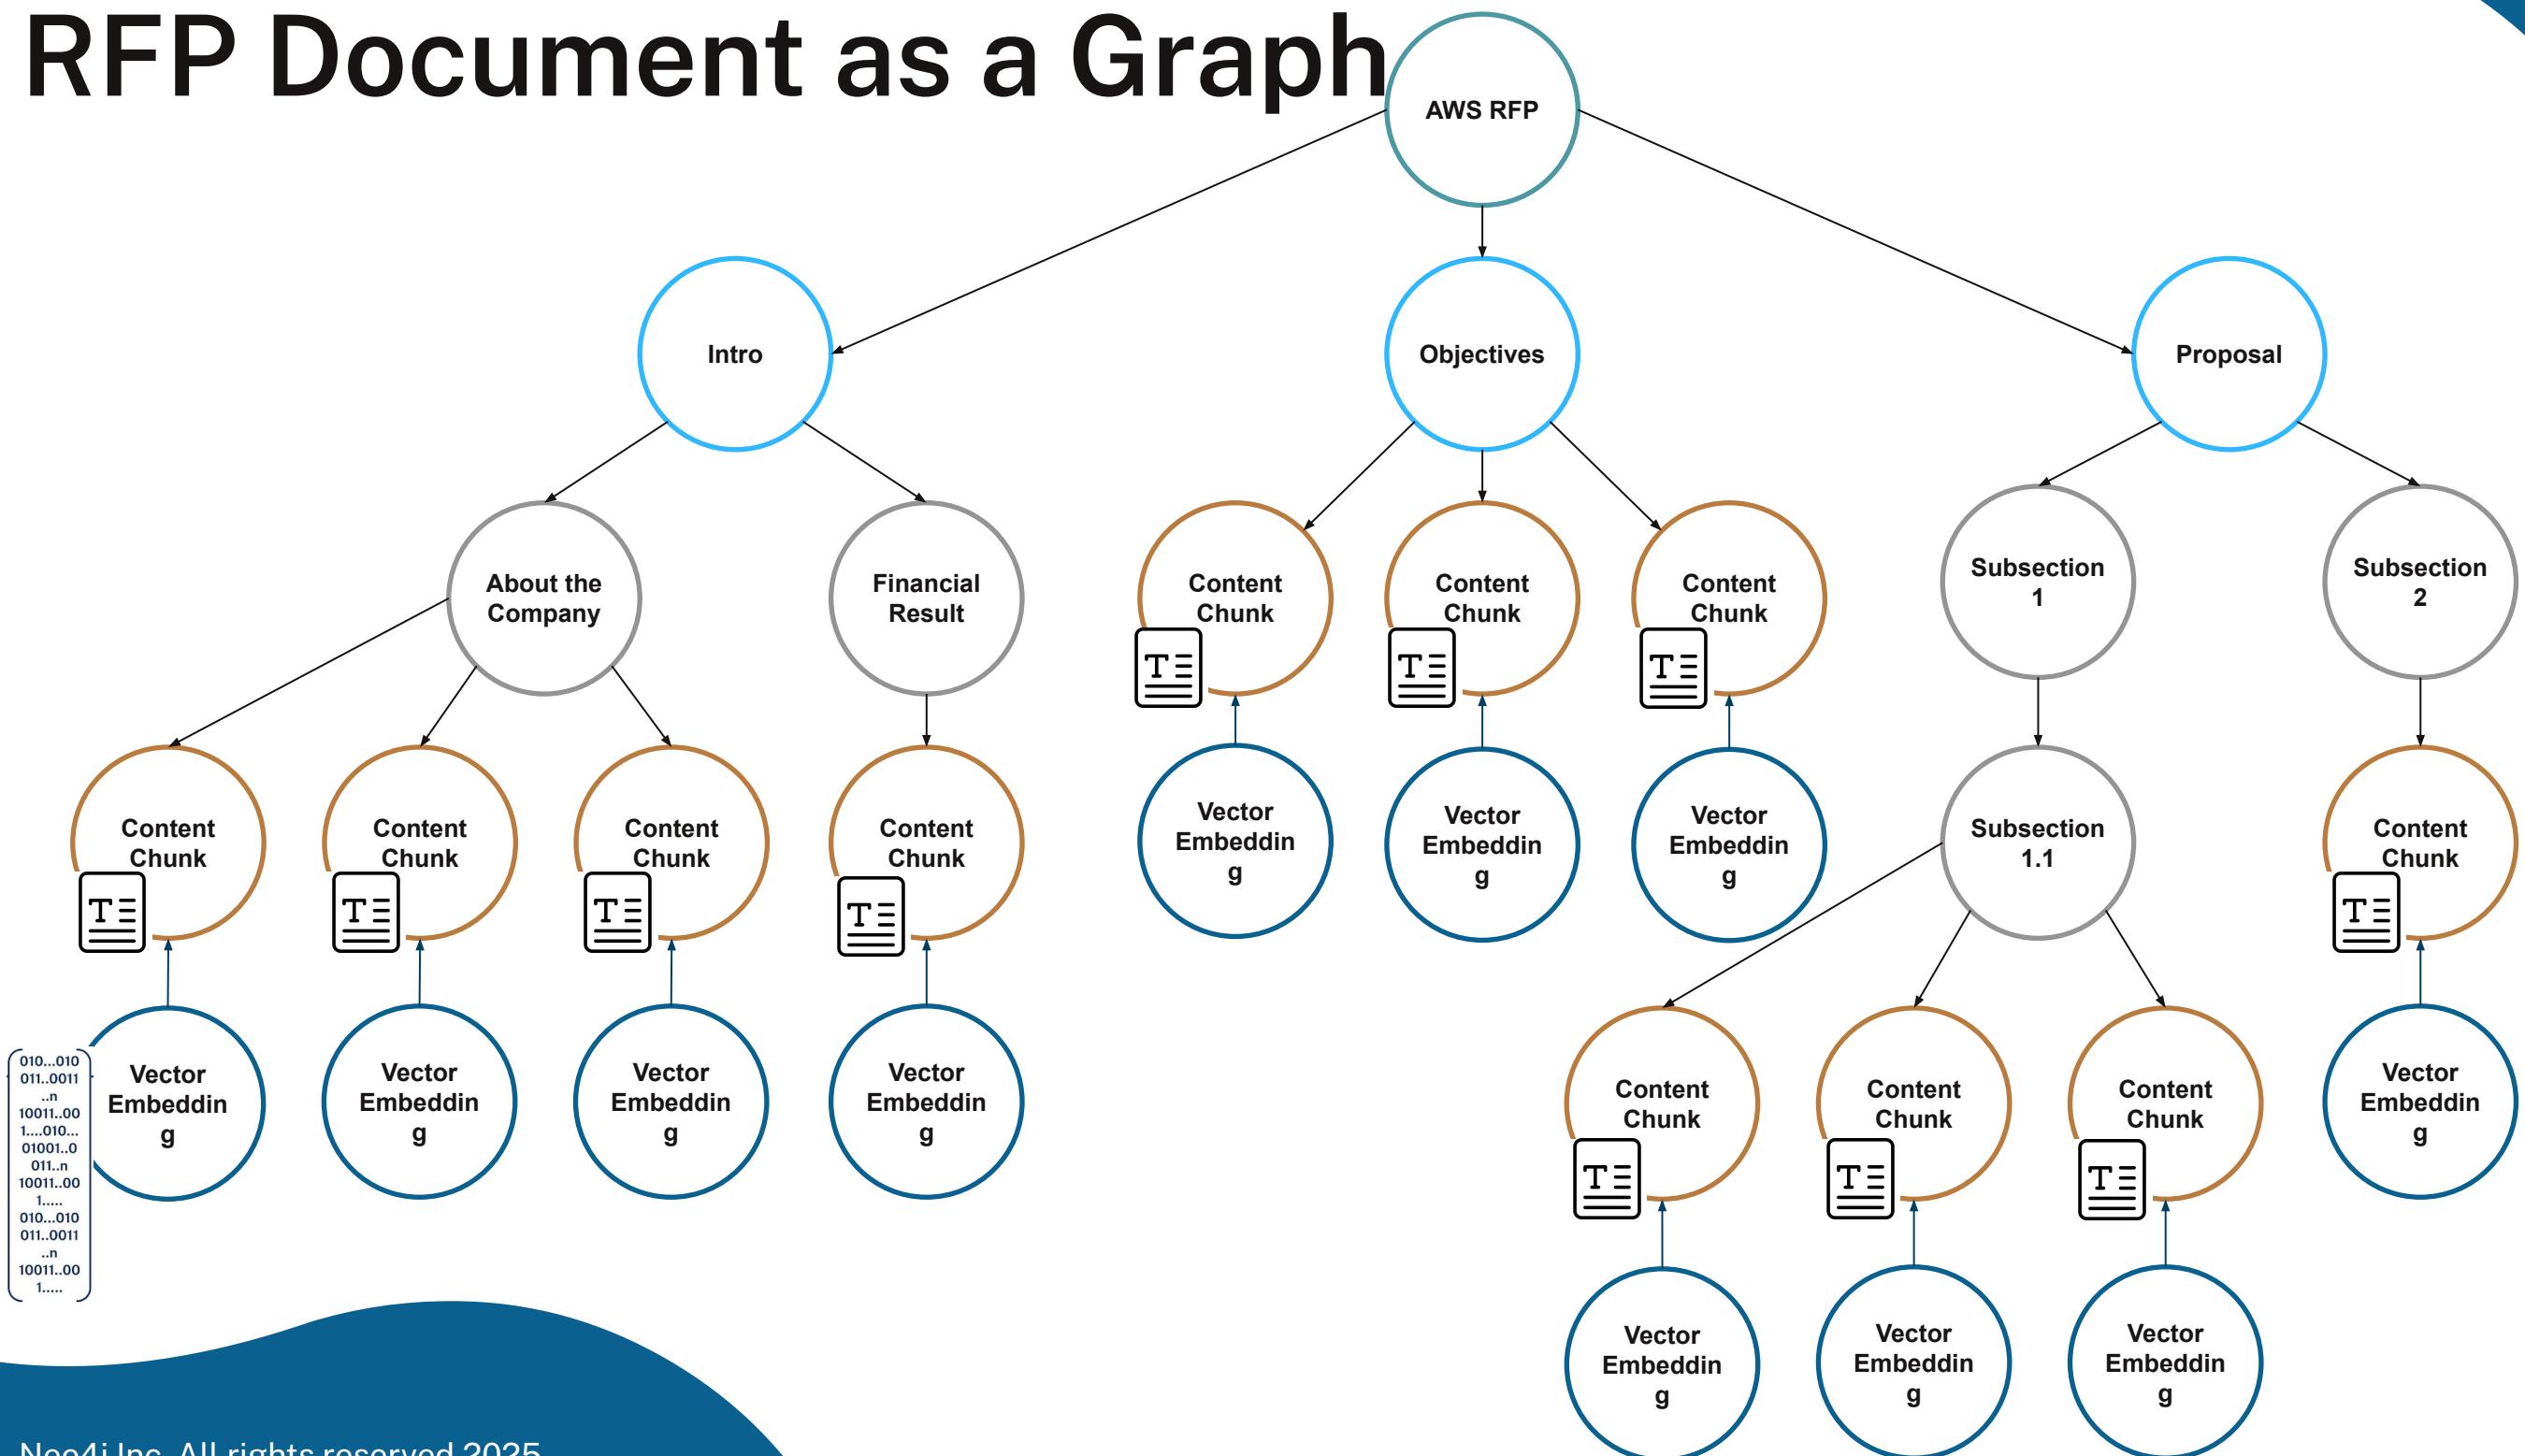
Subsection 1.1

Content

Subsection 2

Content

RFP Document as a Graph

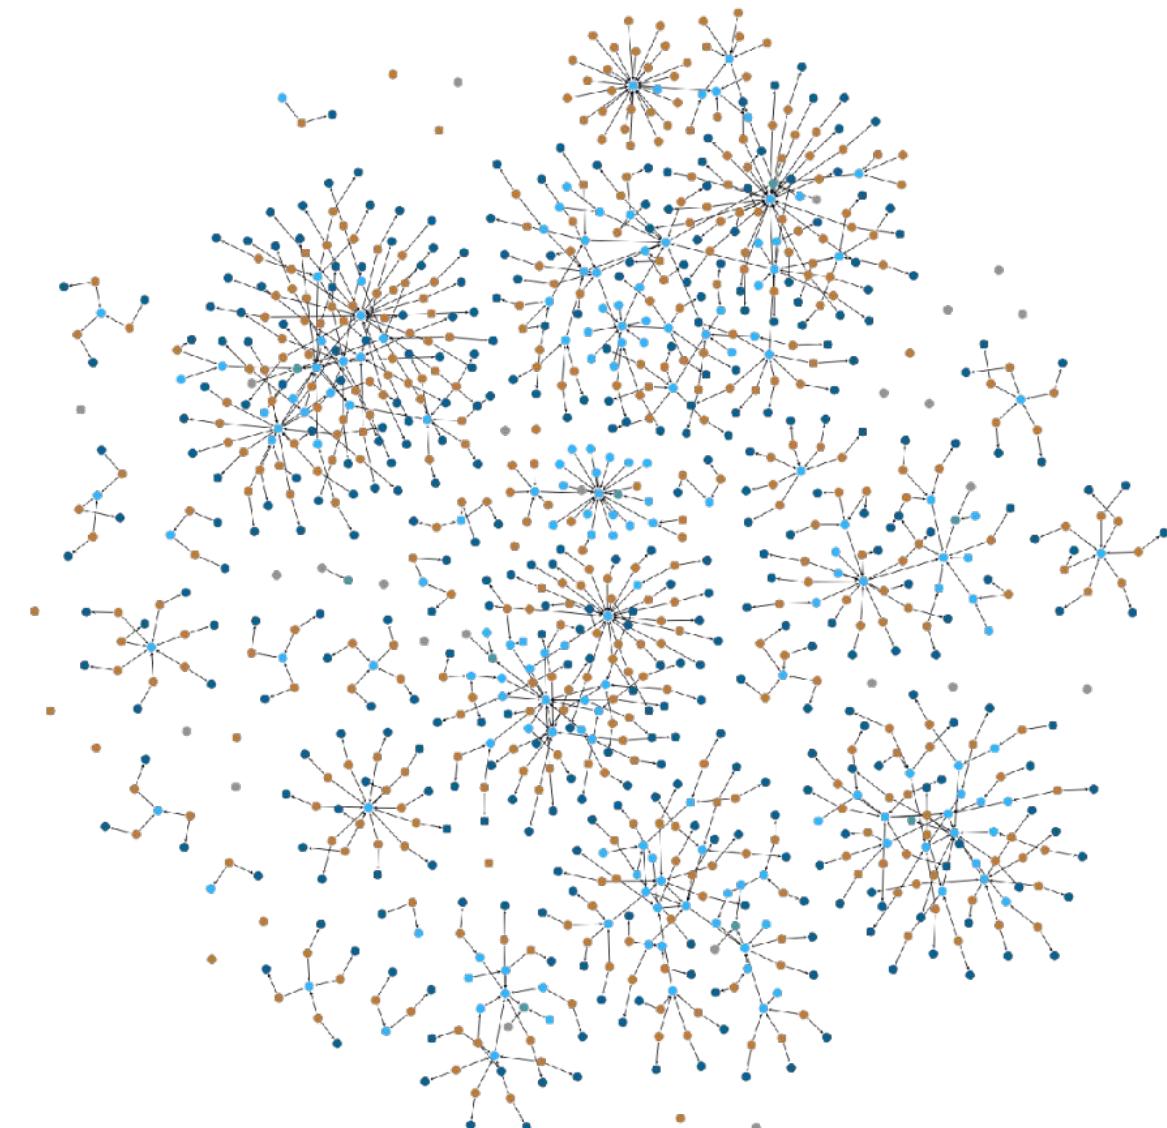
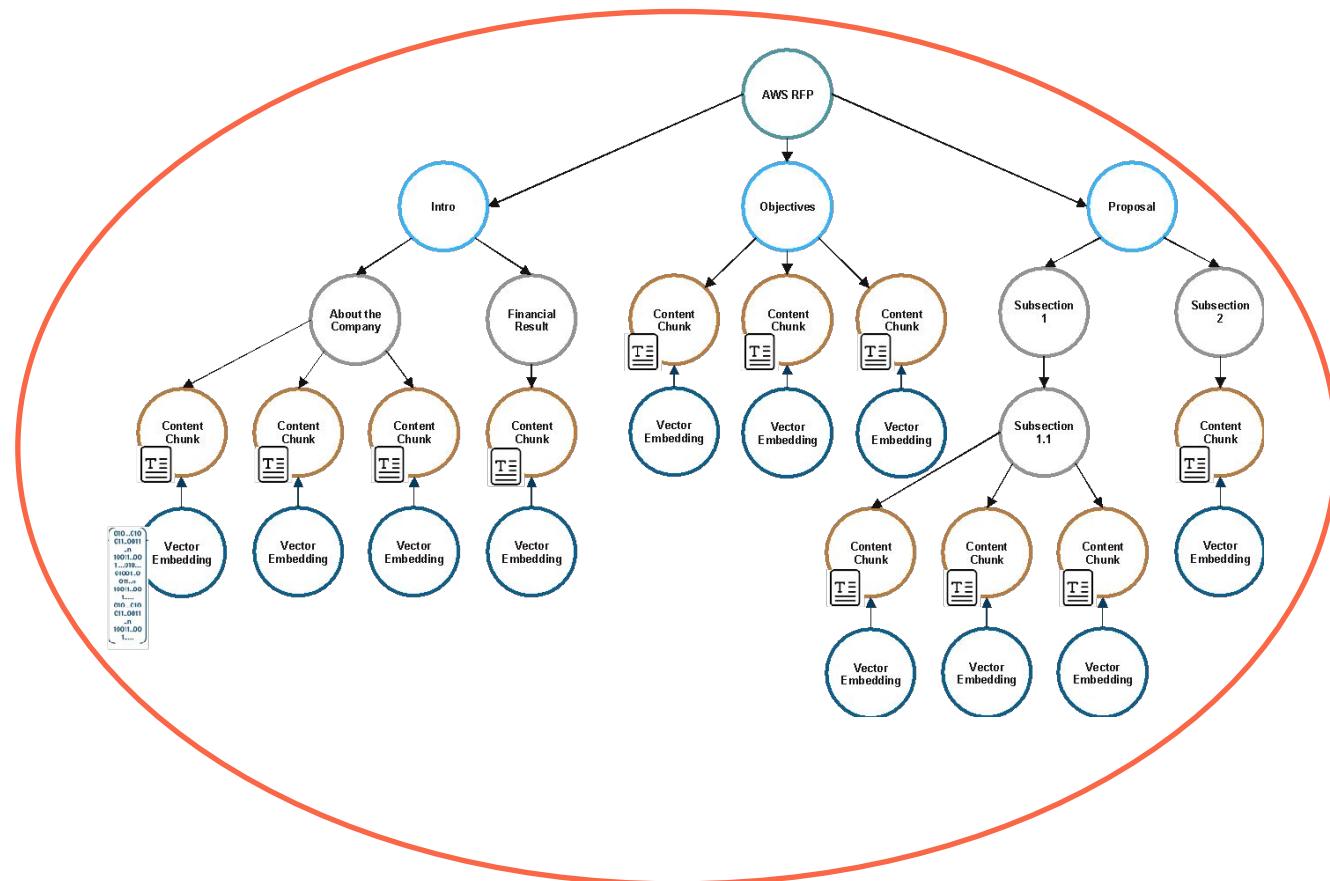


Knowledge Graph as the Knowledge Base

Document in a KG



Knowledge Graph

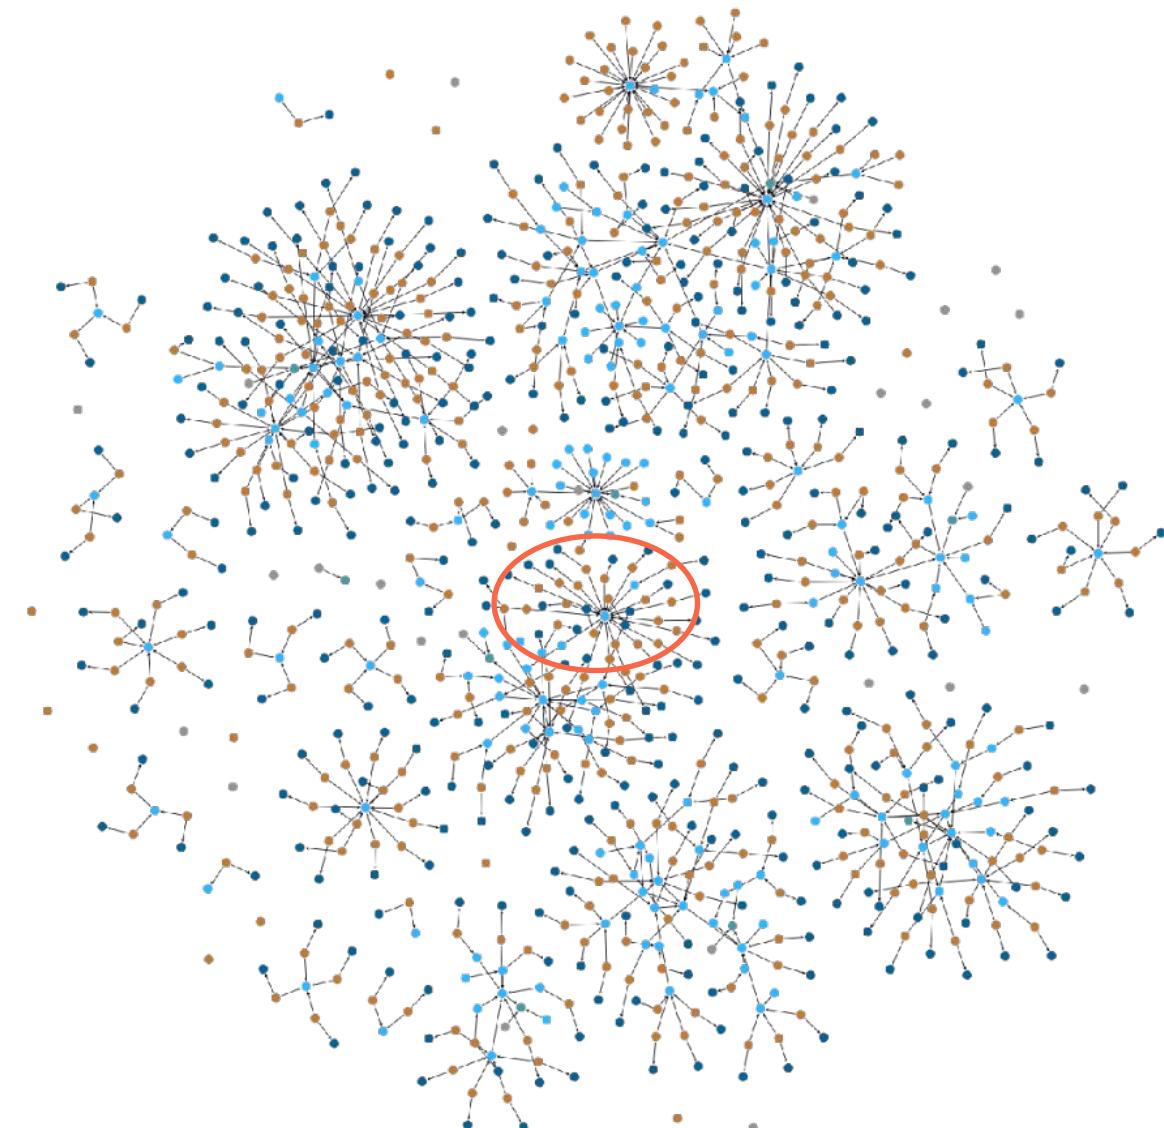
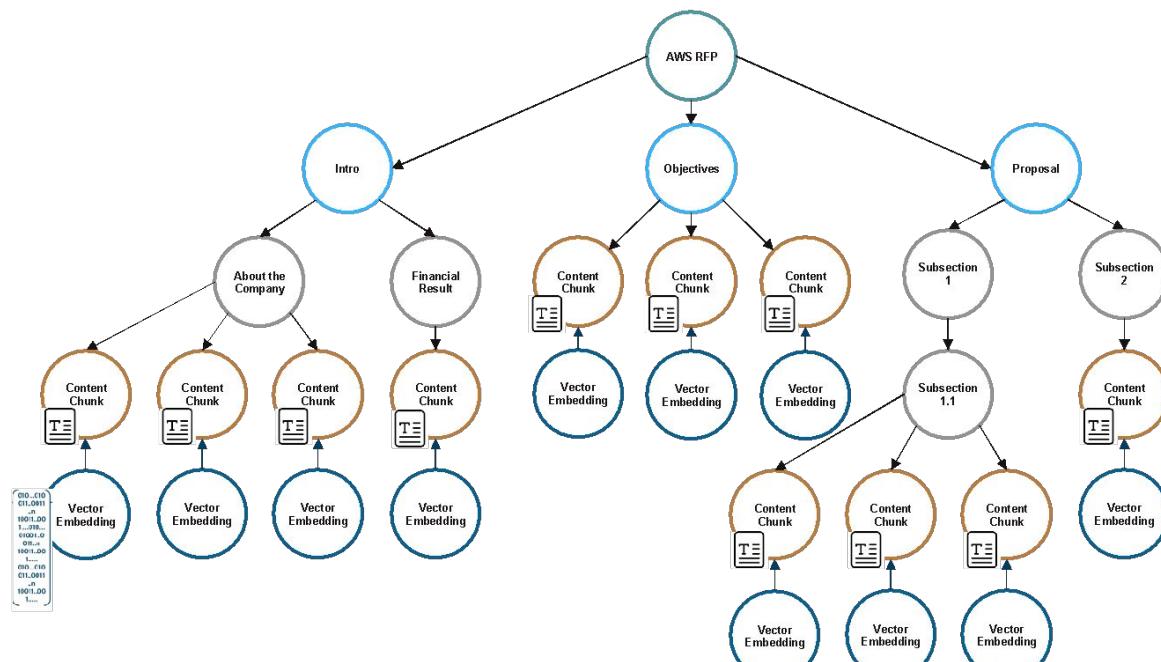


Knowledge Graph as the Knowledge Base

Document in a KG



Knowledge Graph



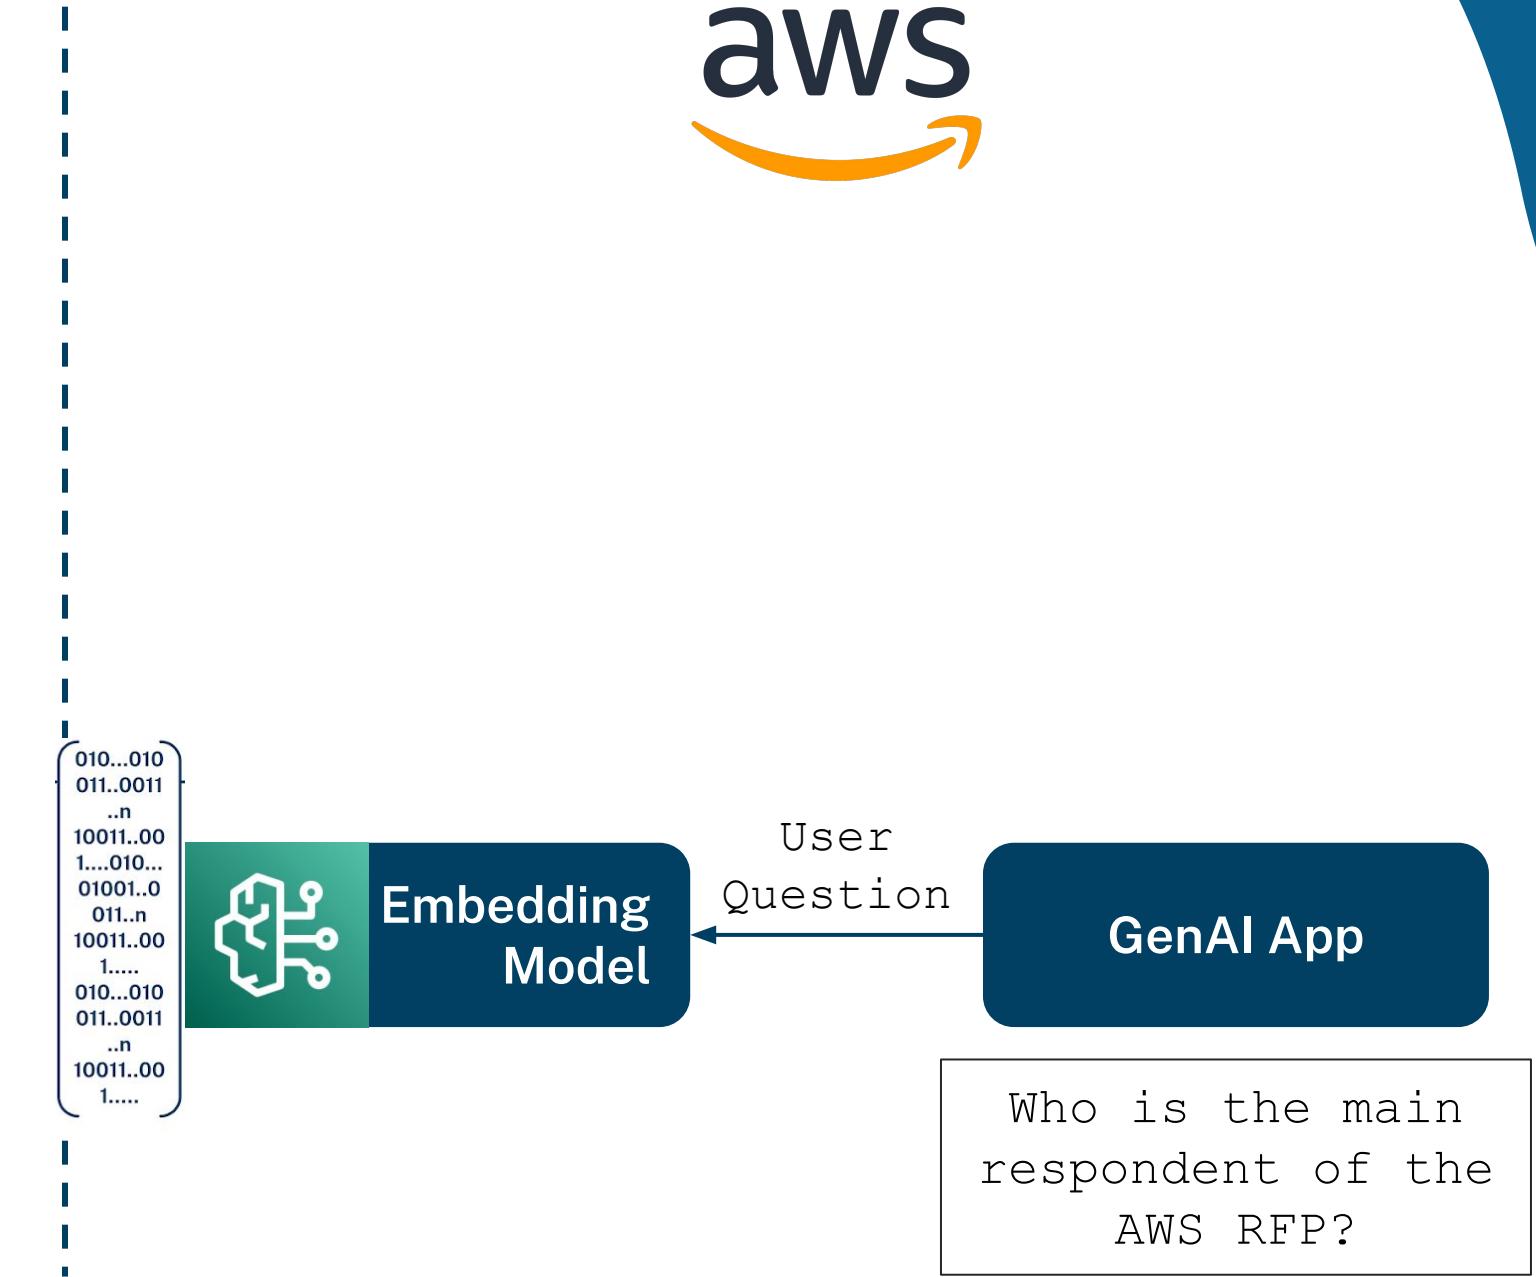


Accurate, Contextual and Explainable

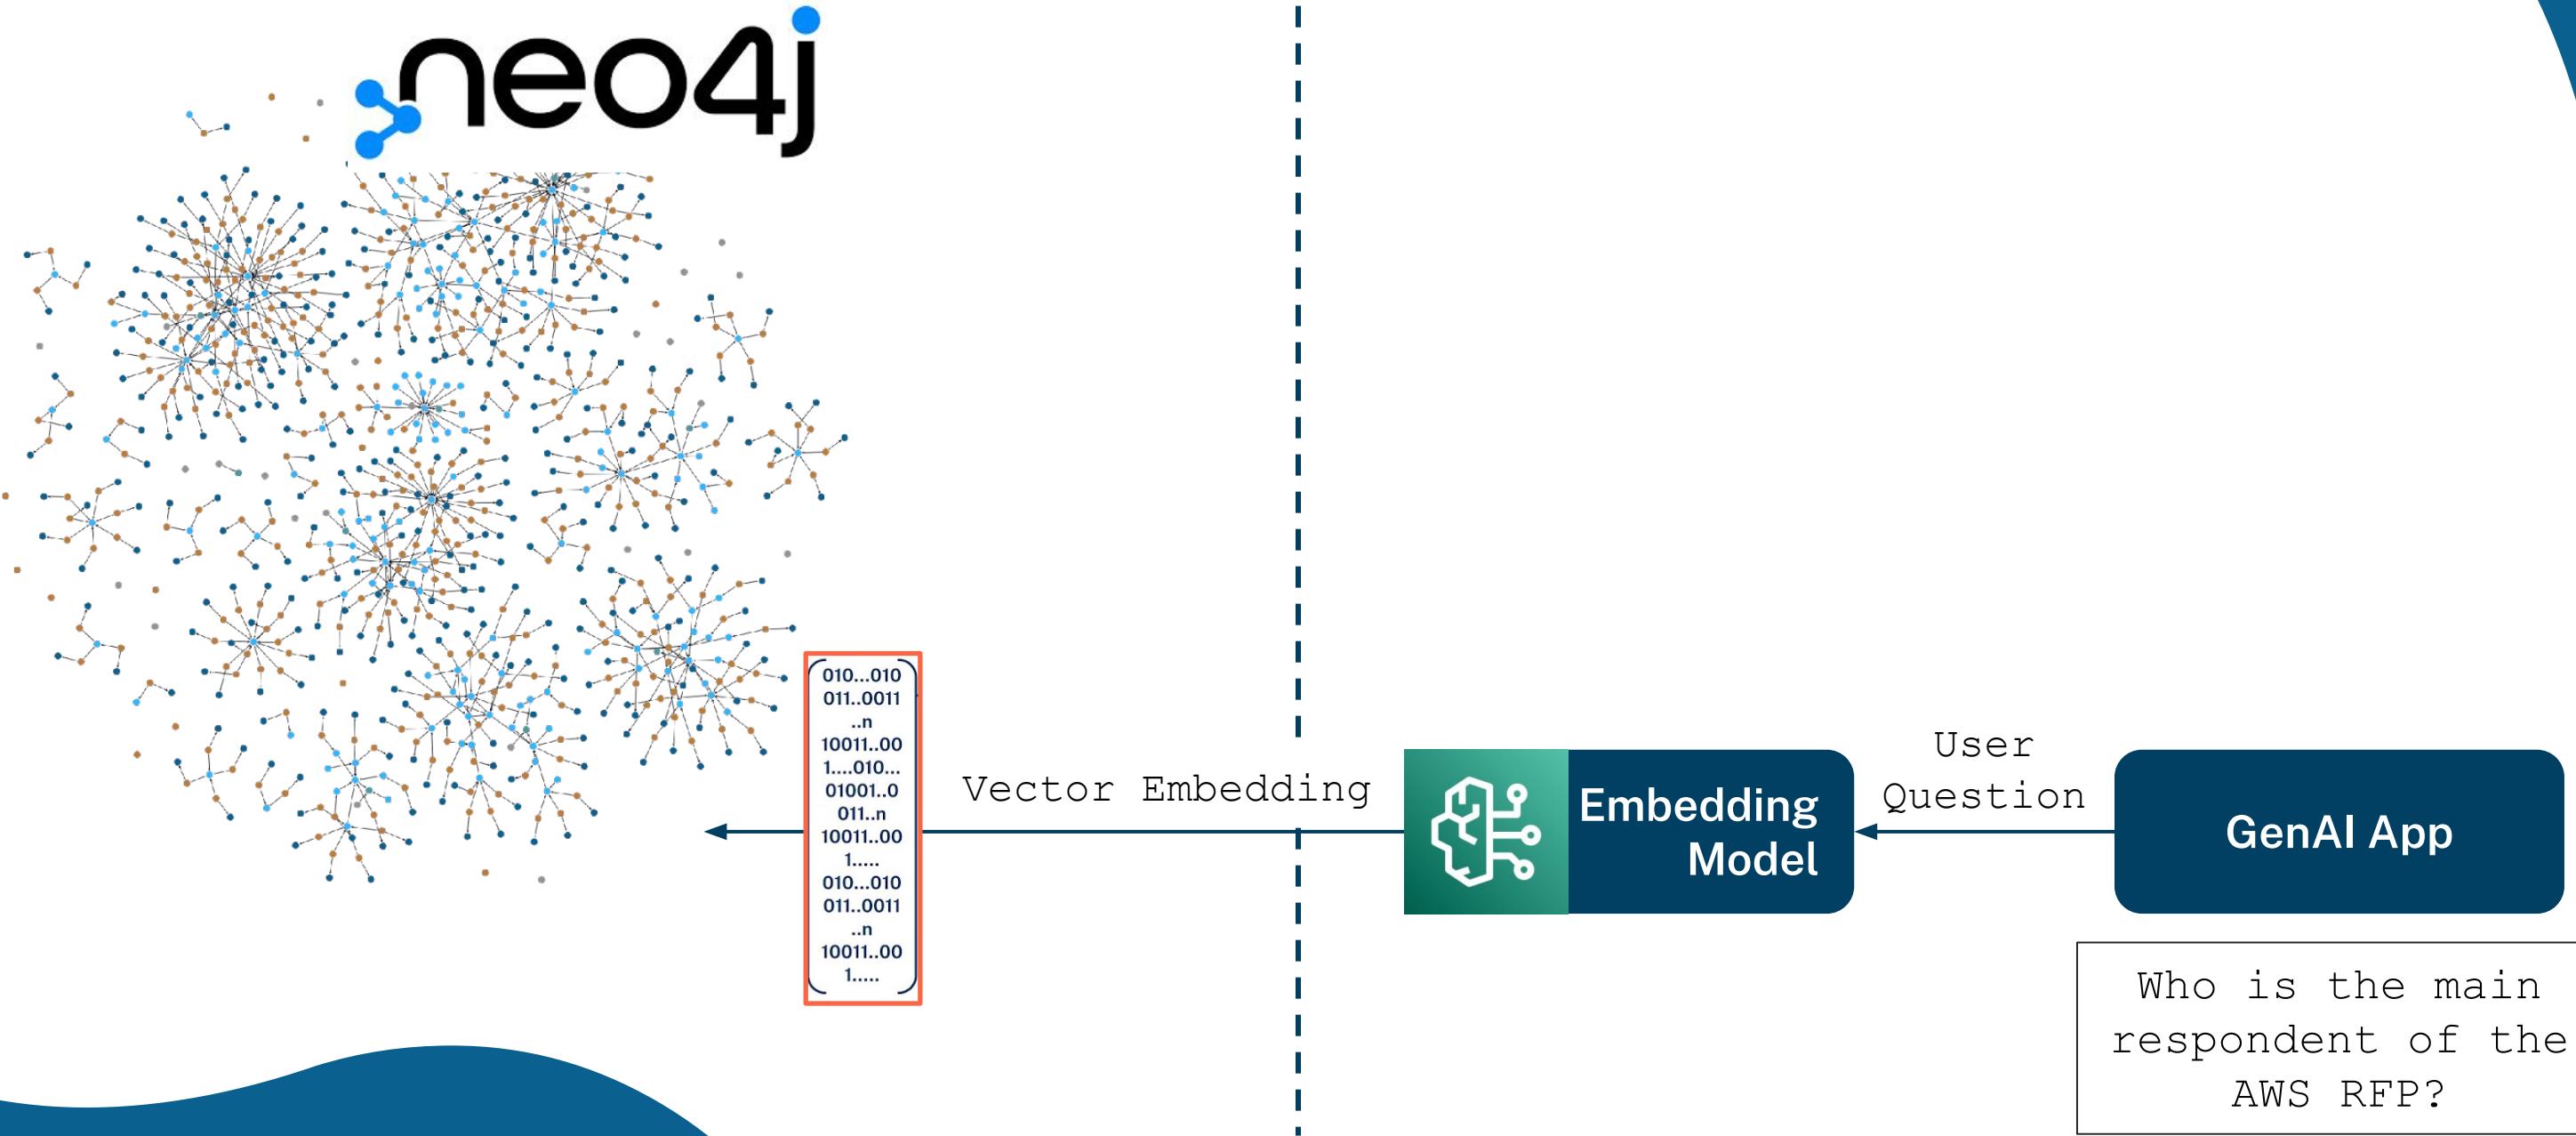
Who is the main respondent
of the AWS RFP?

GenAI App

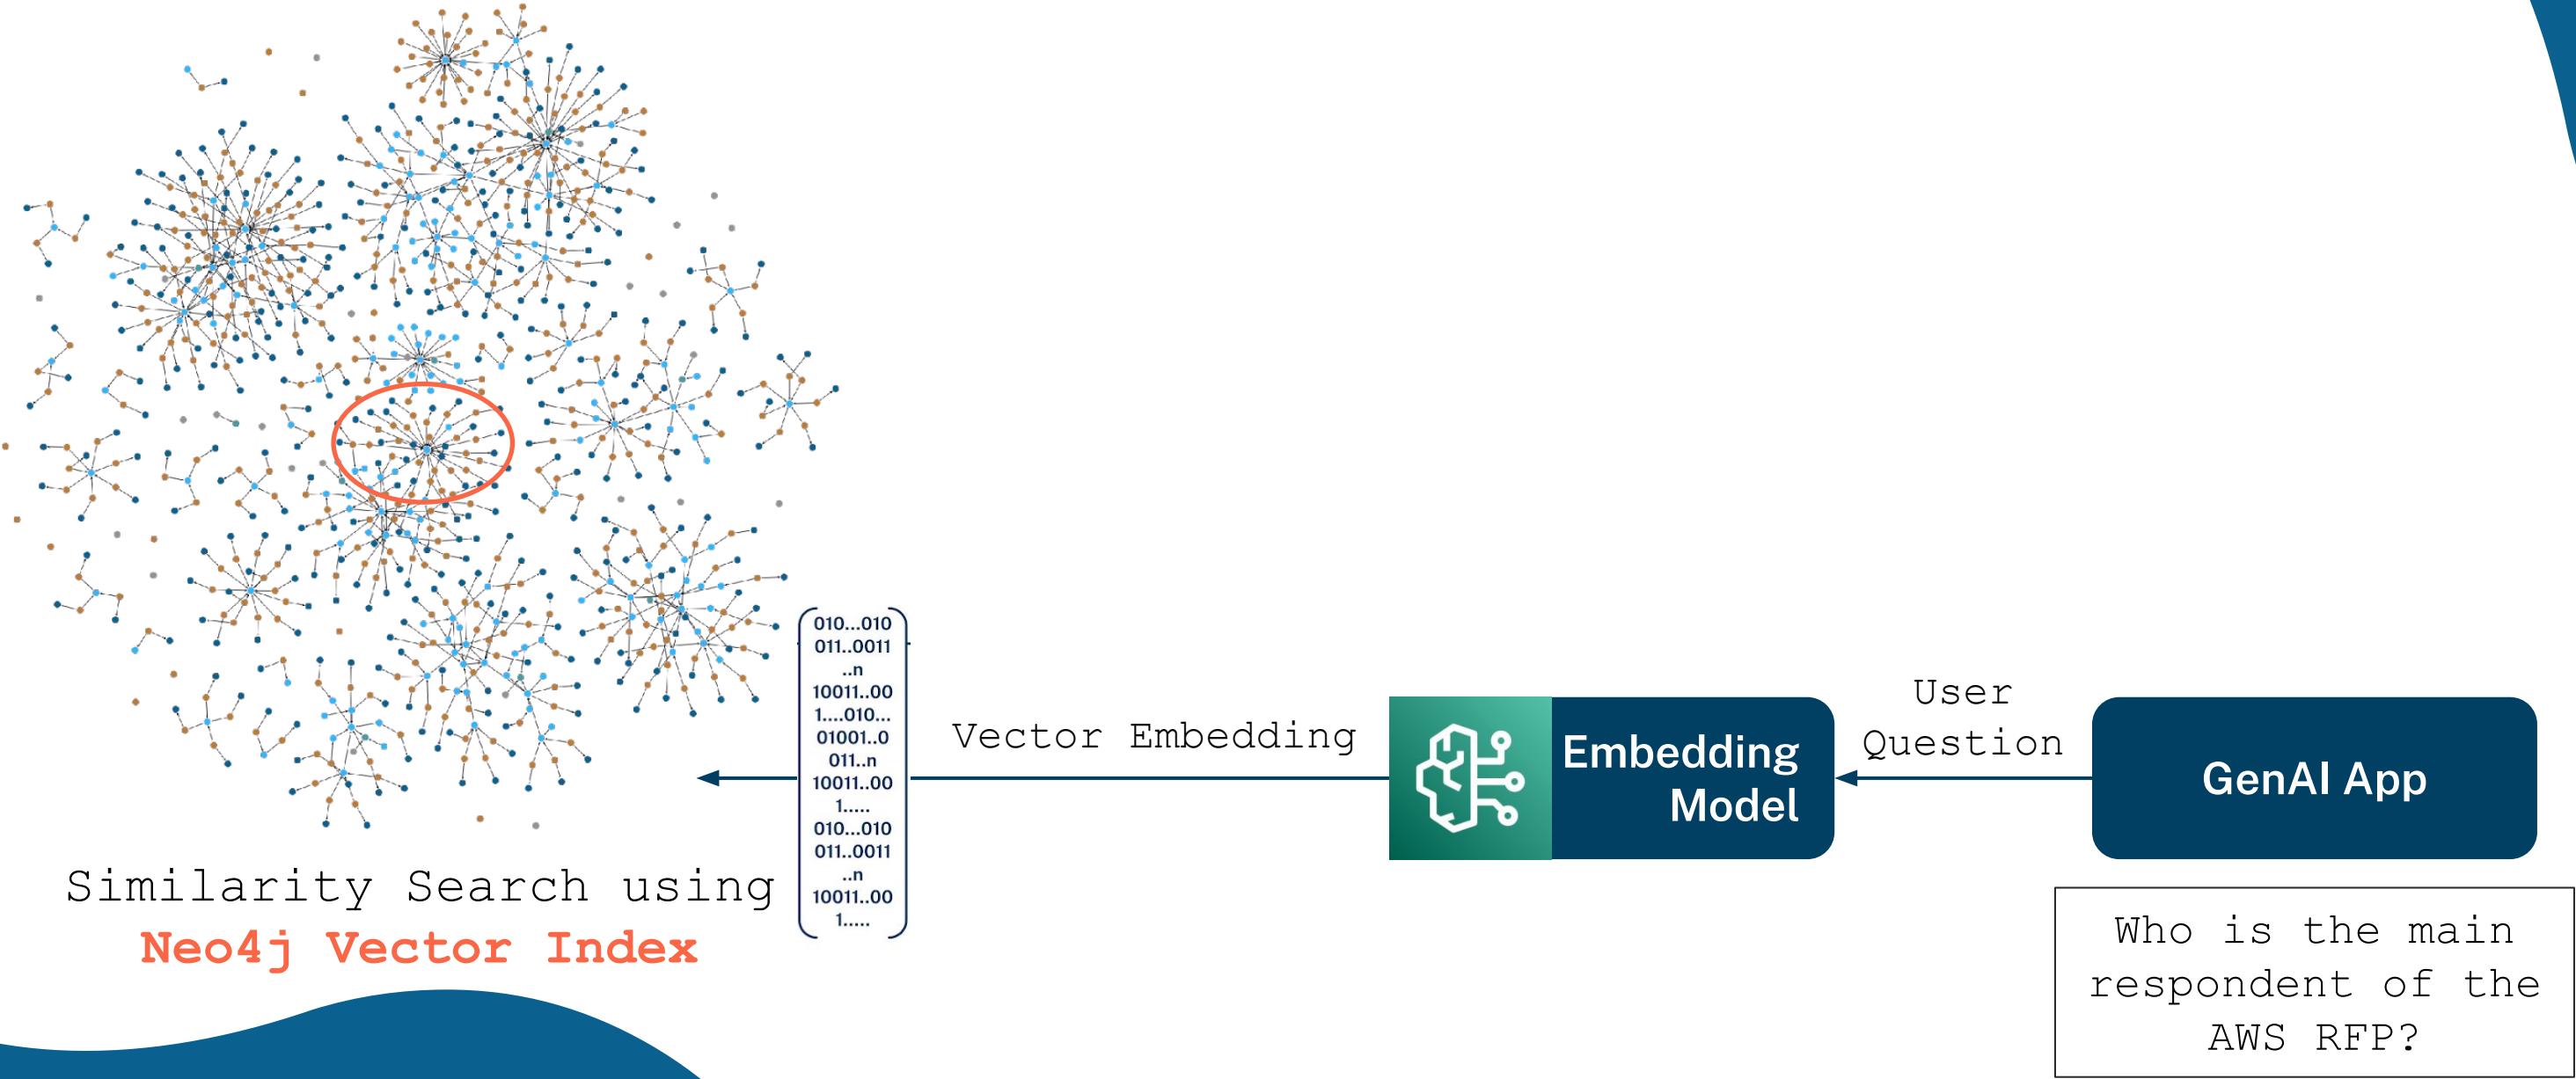
• Accurate, Contextual and Explainable



neo4j Accurate, Contextual and Explainable

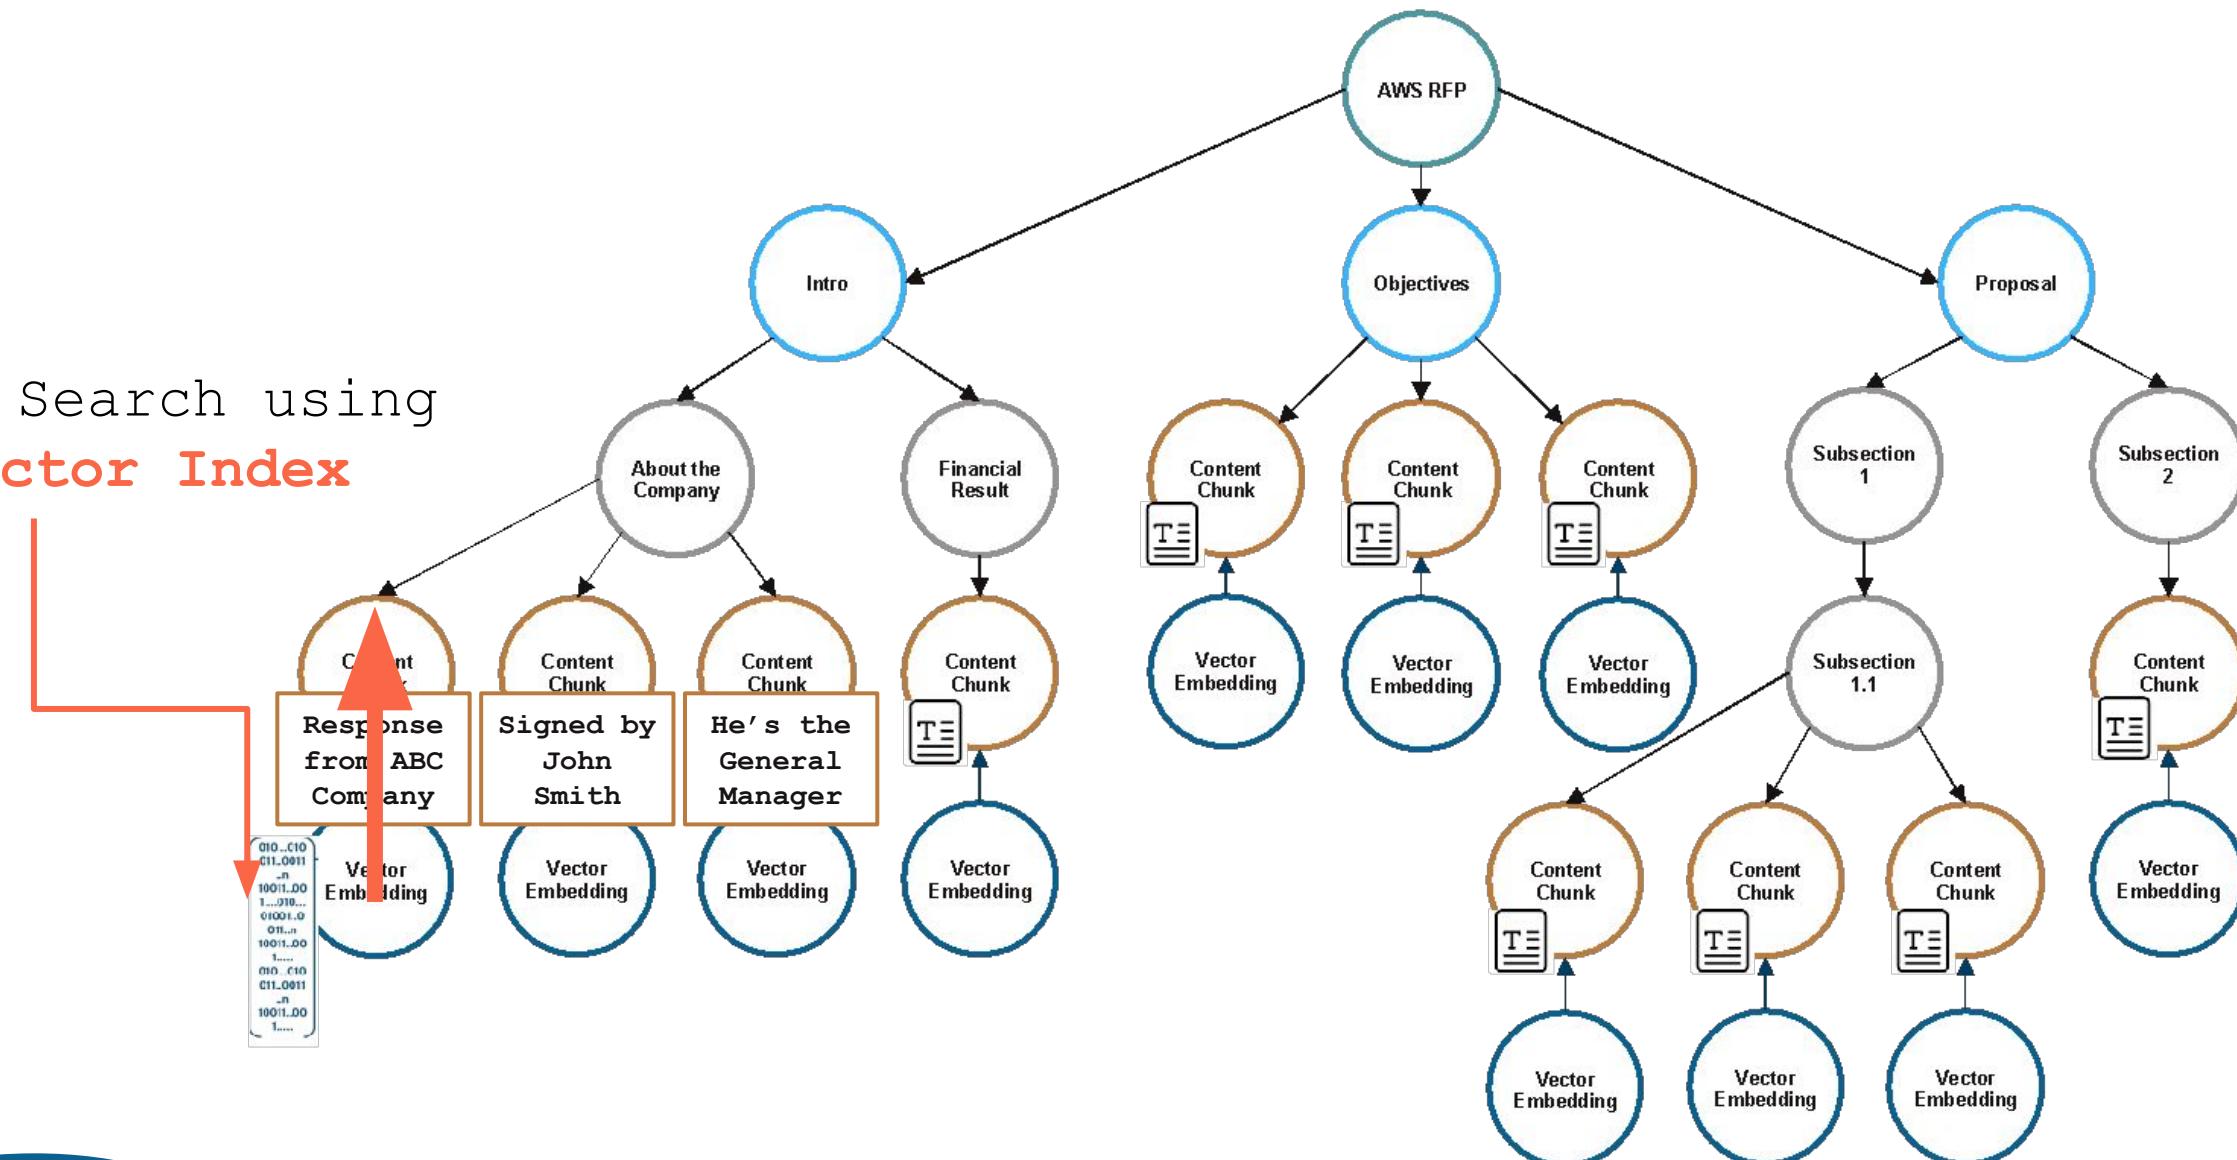


• Accurate, Contextual and Explainable



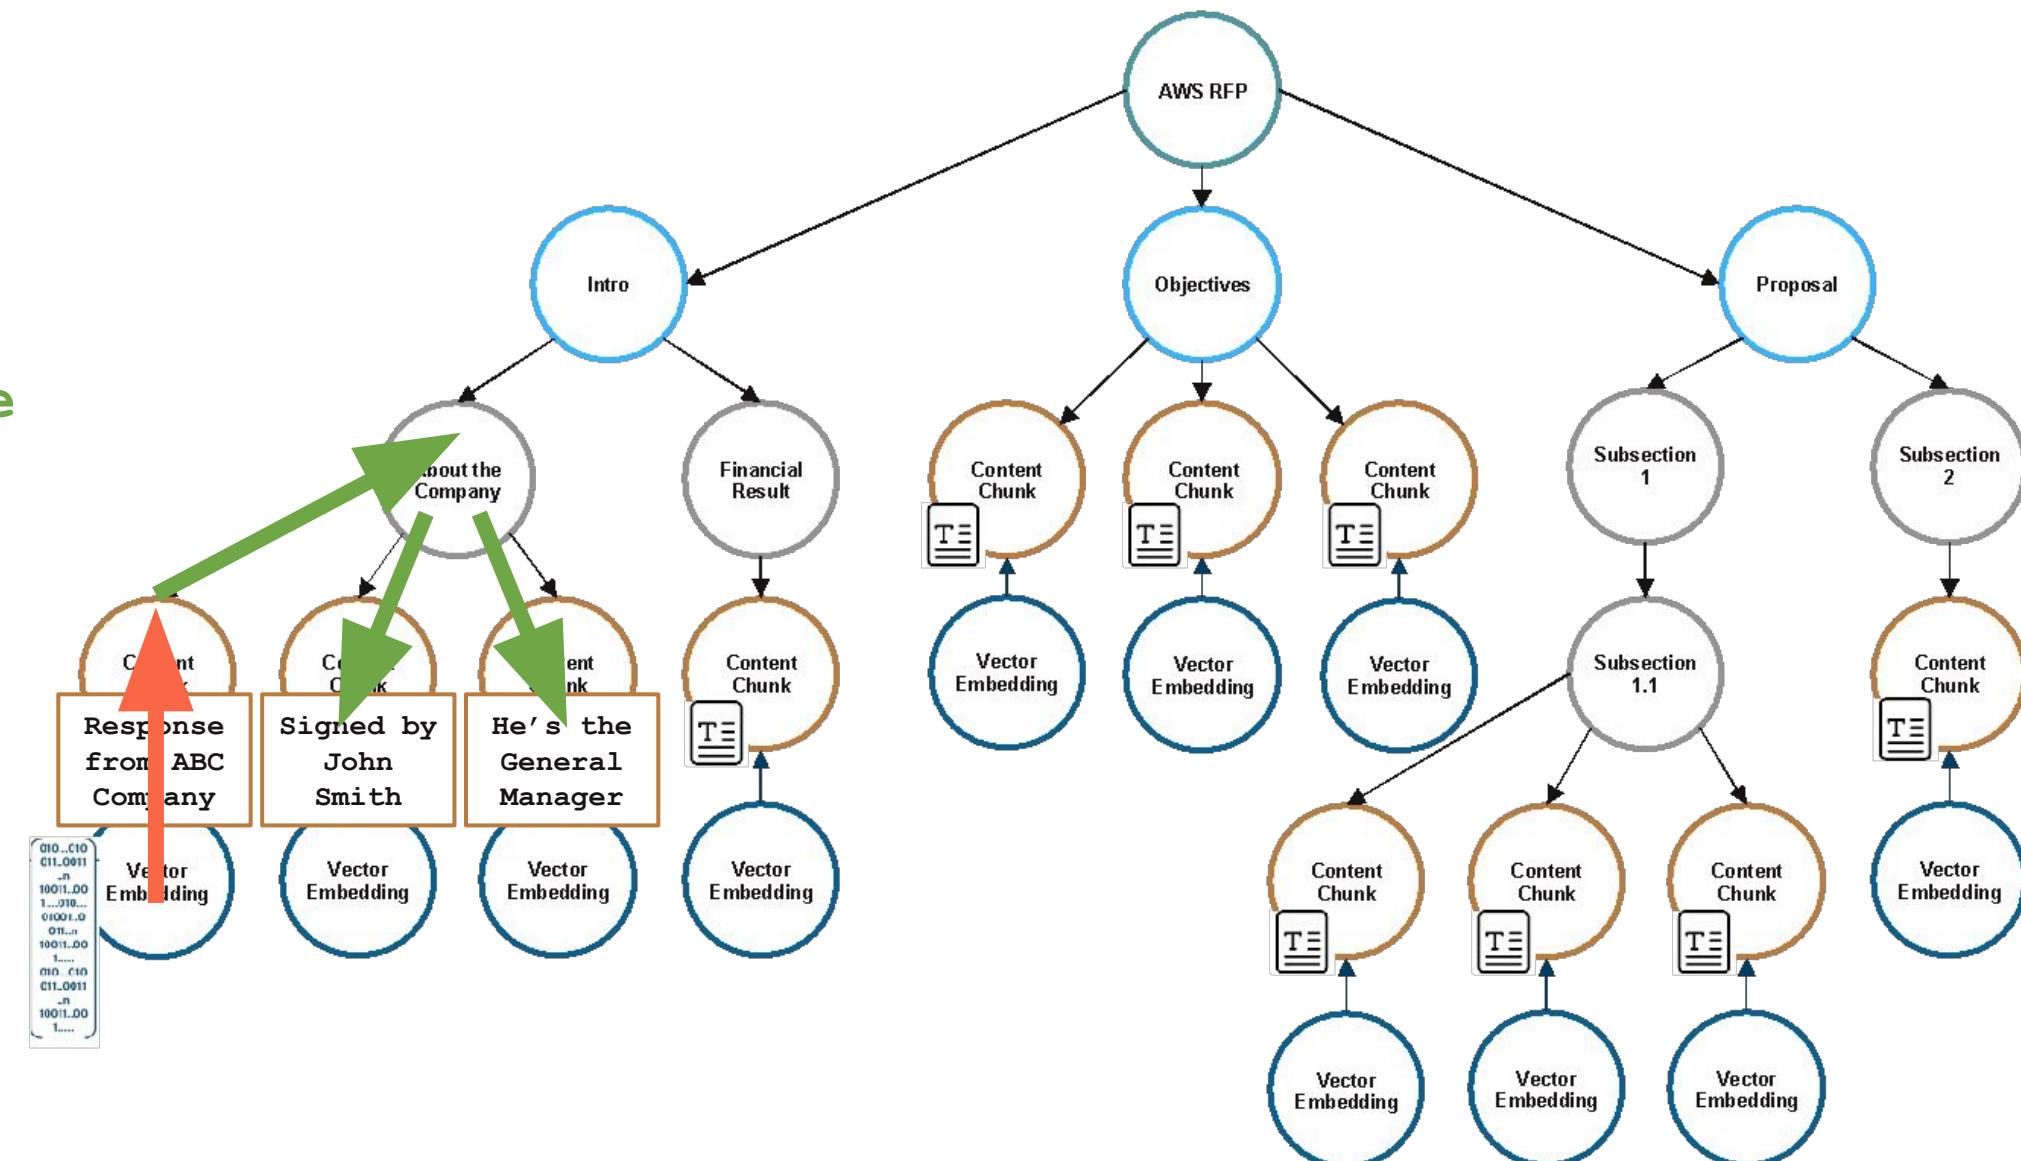
• Accurate, Contextual and Explainable

Similarity Search using
Neo4j Vector Index



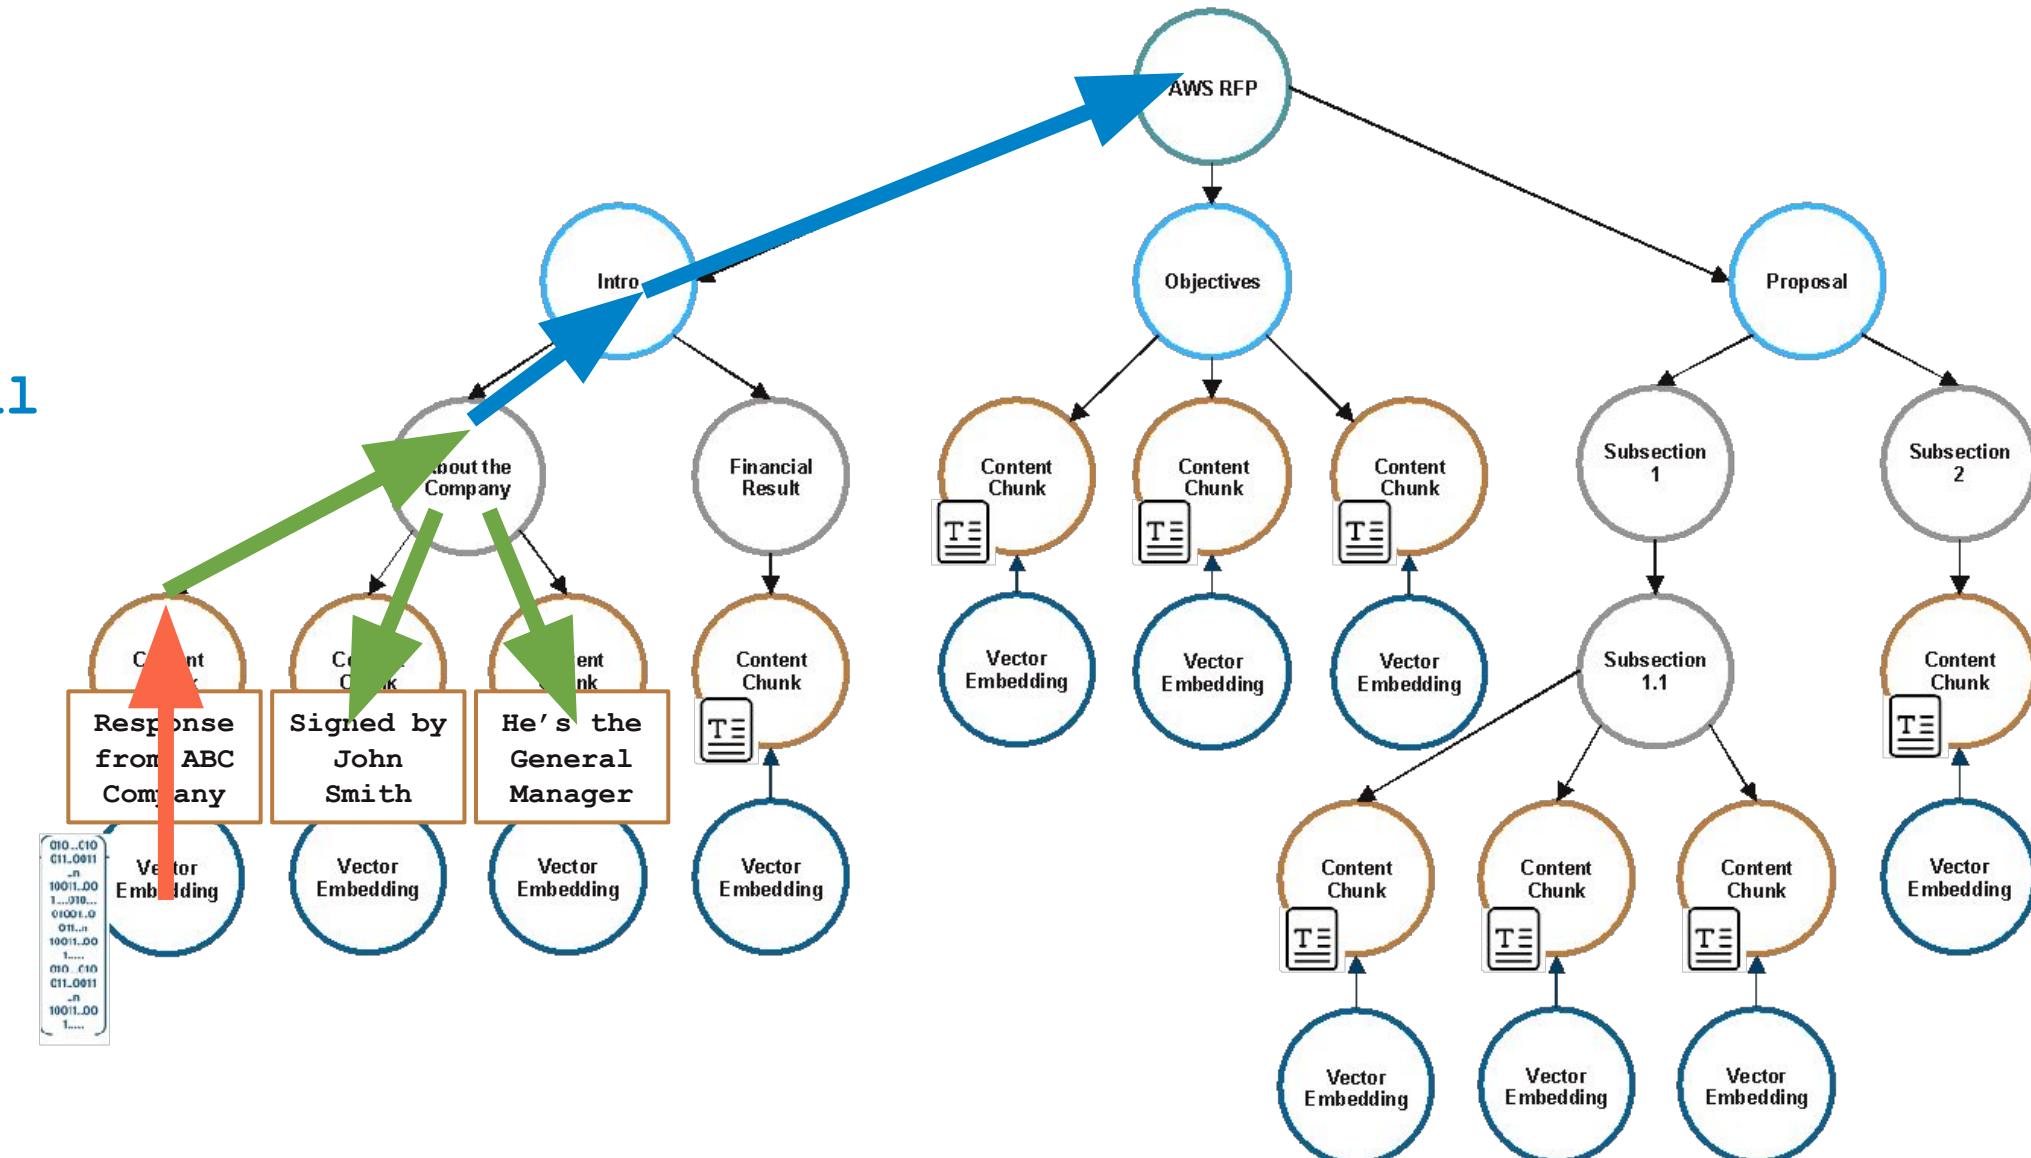
• Accurate, Contextual and Explainable

Contextual Knowledge
Retrieval within
Neo4j KG



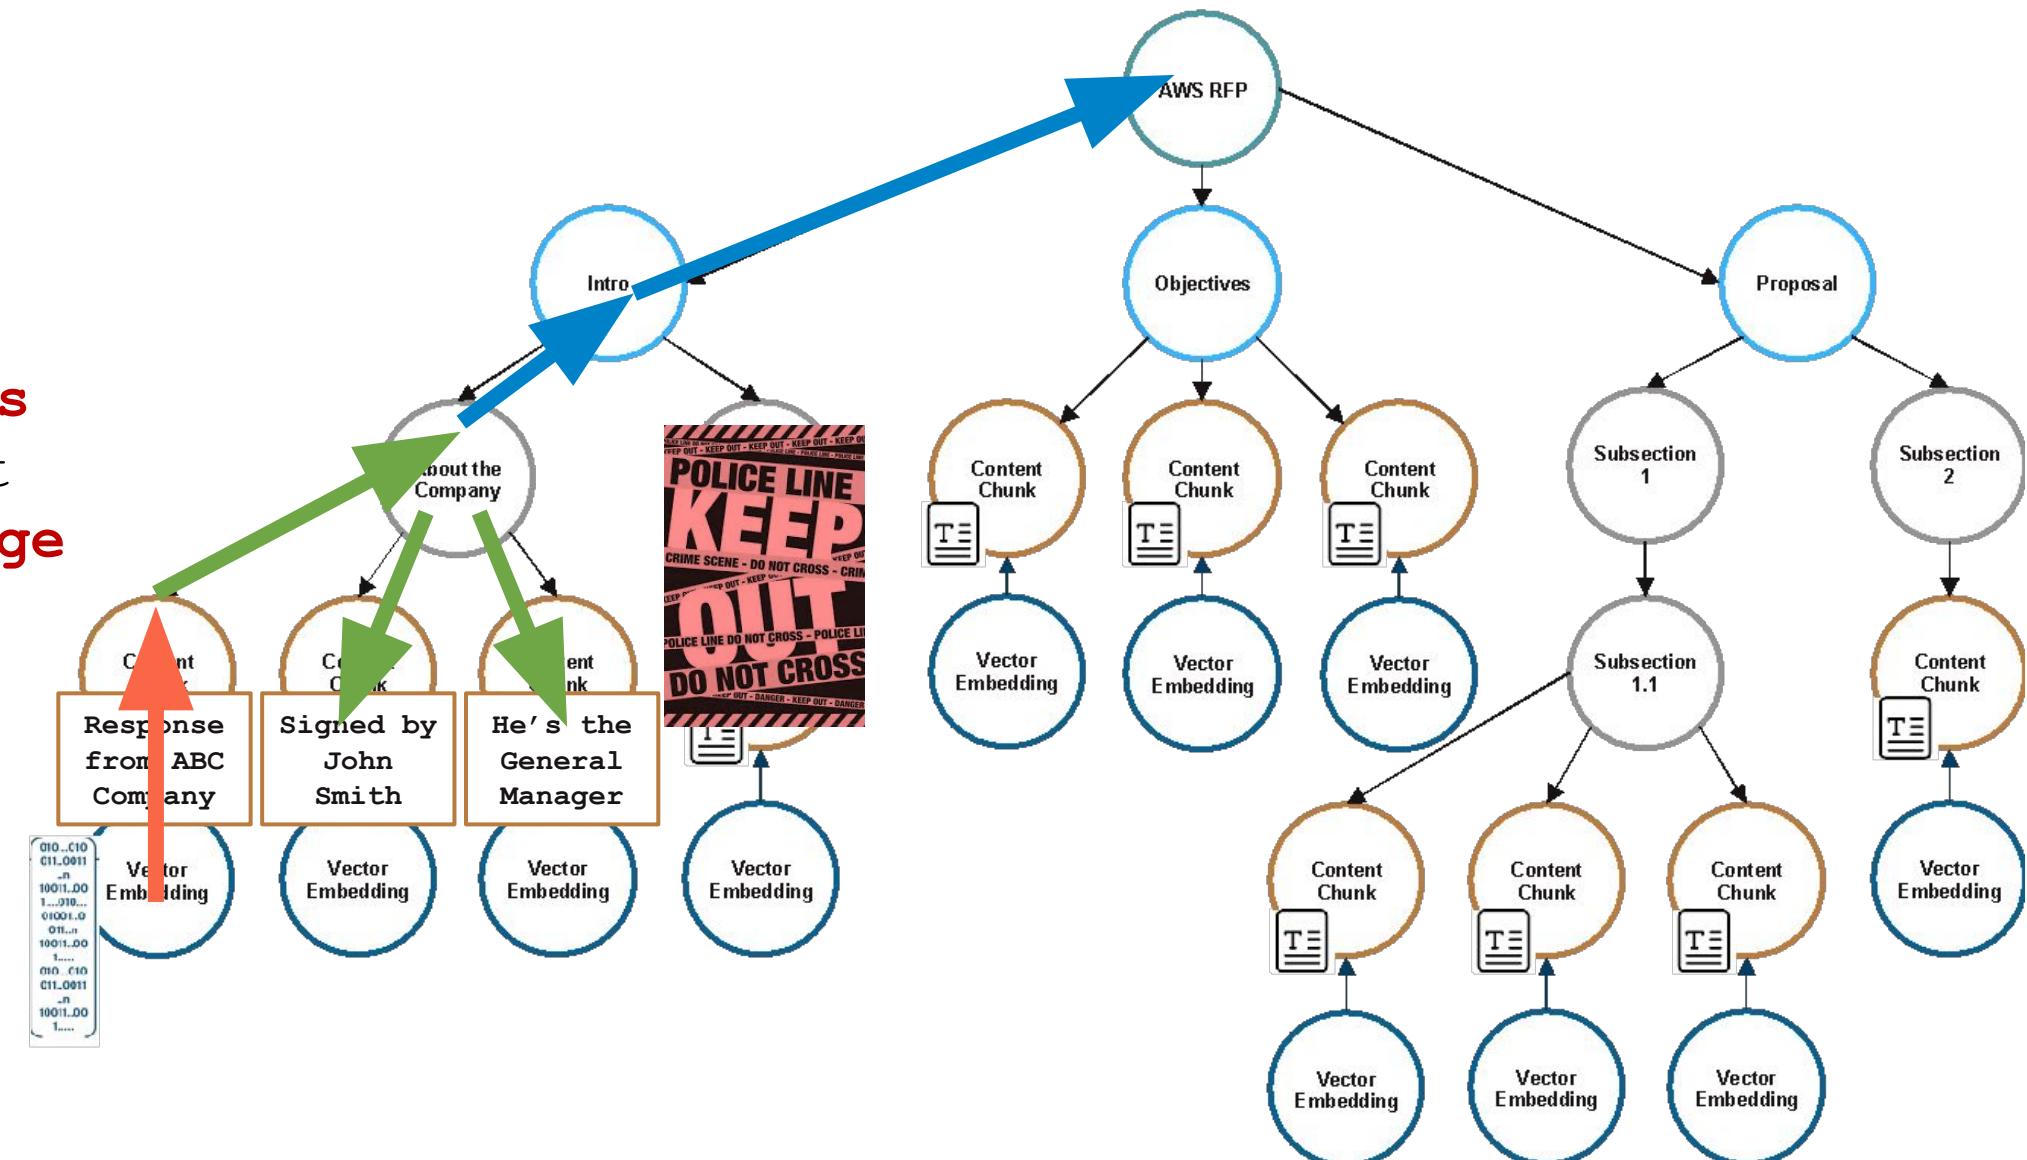
• Accurate, Contextual and Explainable

Knowledge Retrieval
to aid in
Explainability

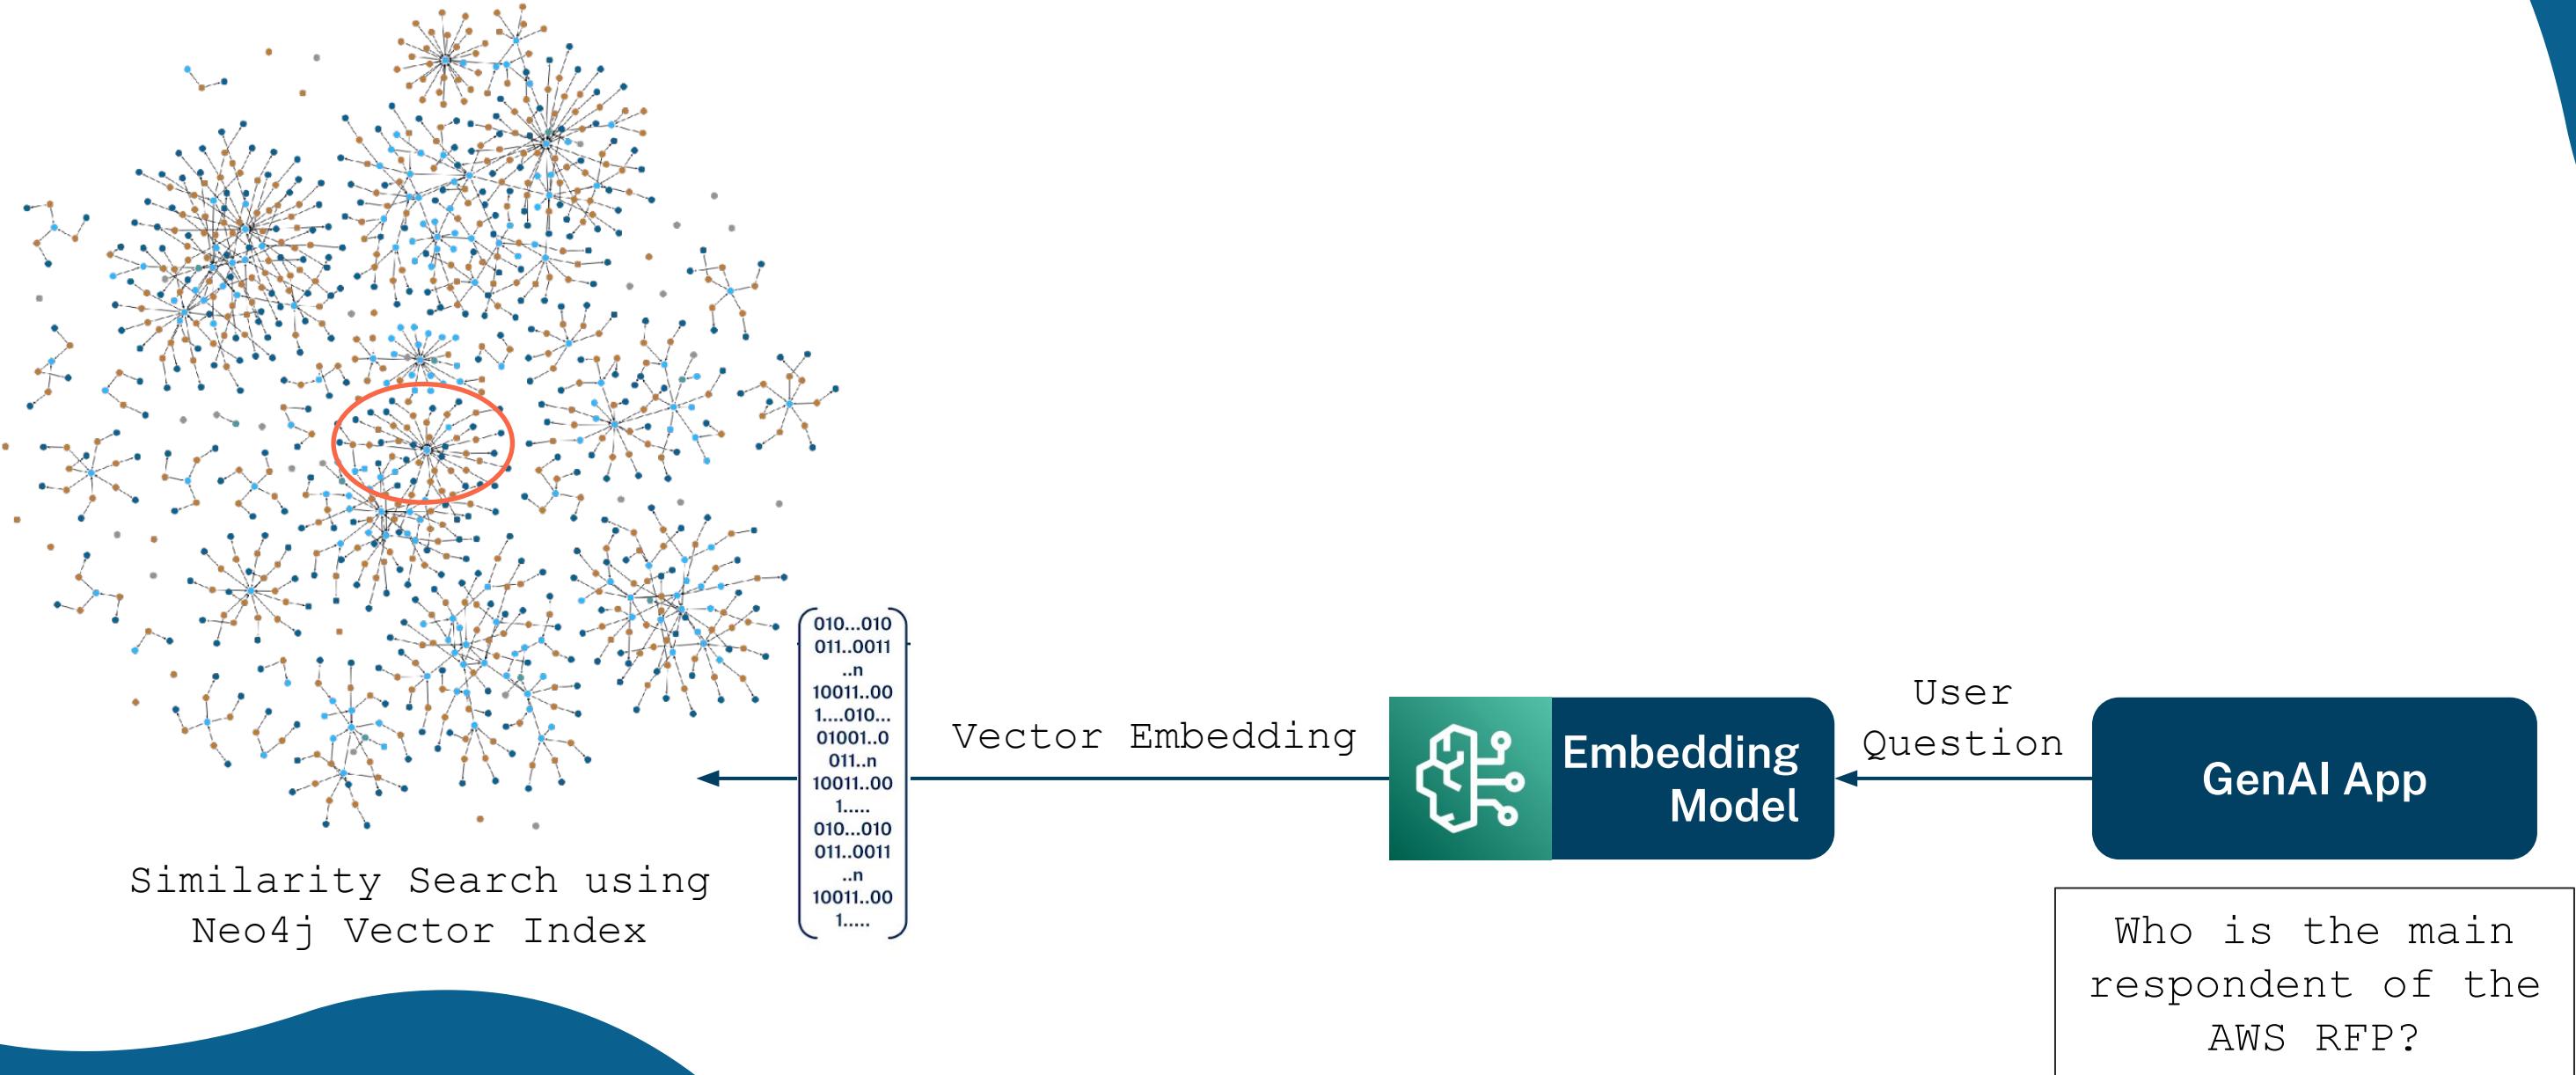


• Accurate, Contextual and Explainable

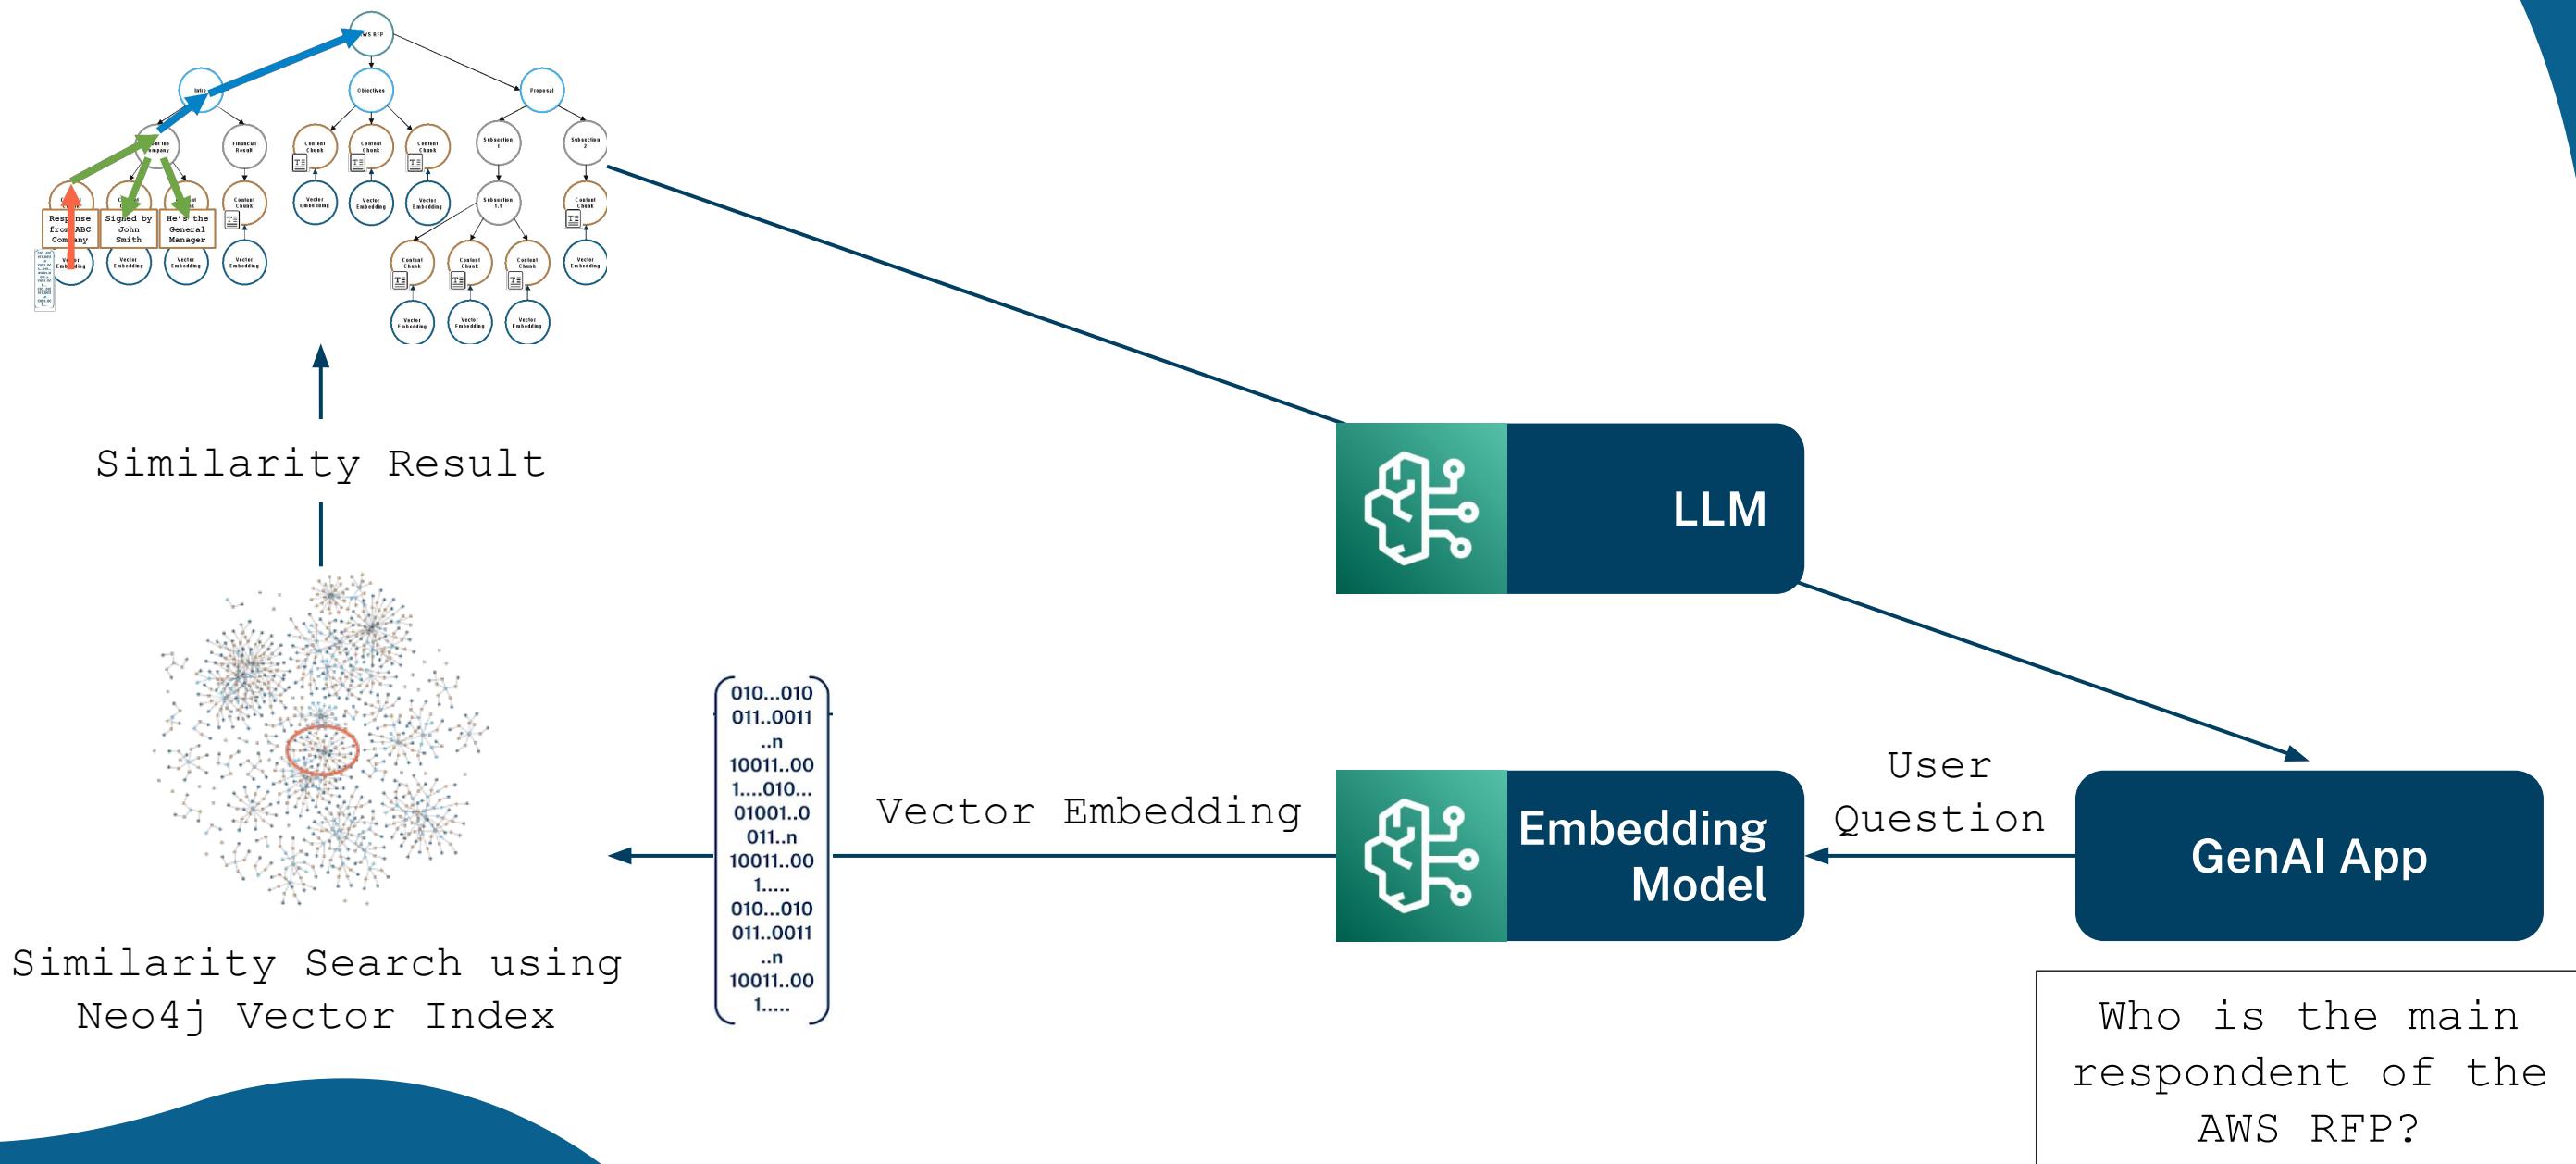
Fine Grained Access Control to prevent unwarranted Knowledge Retrieval



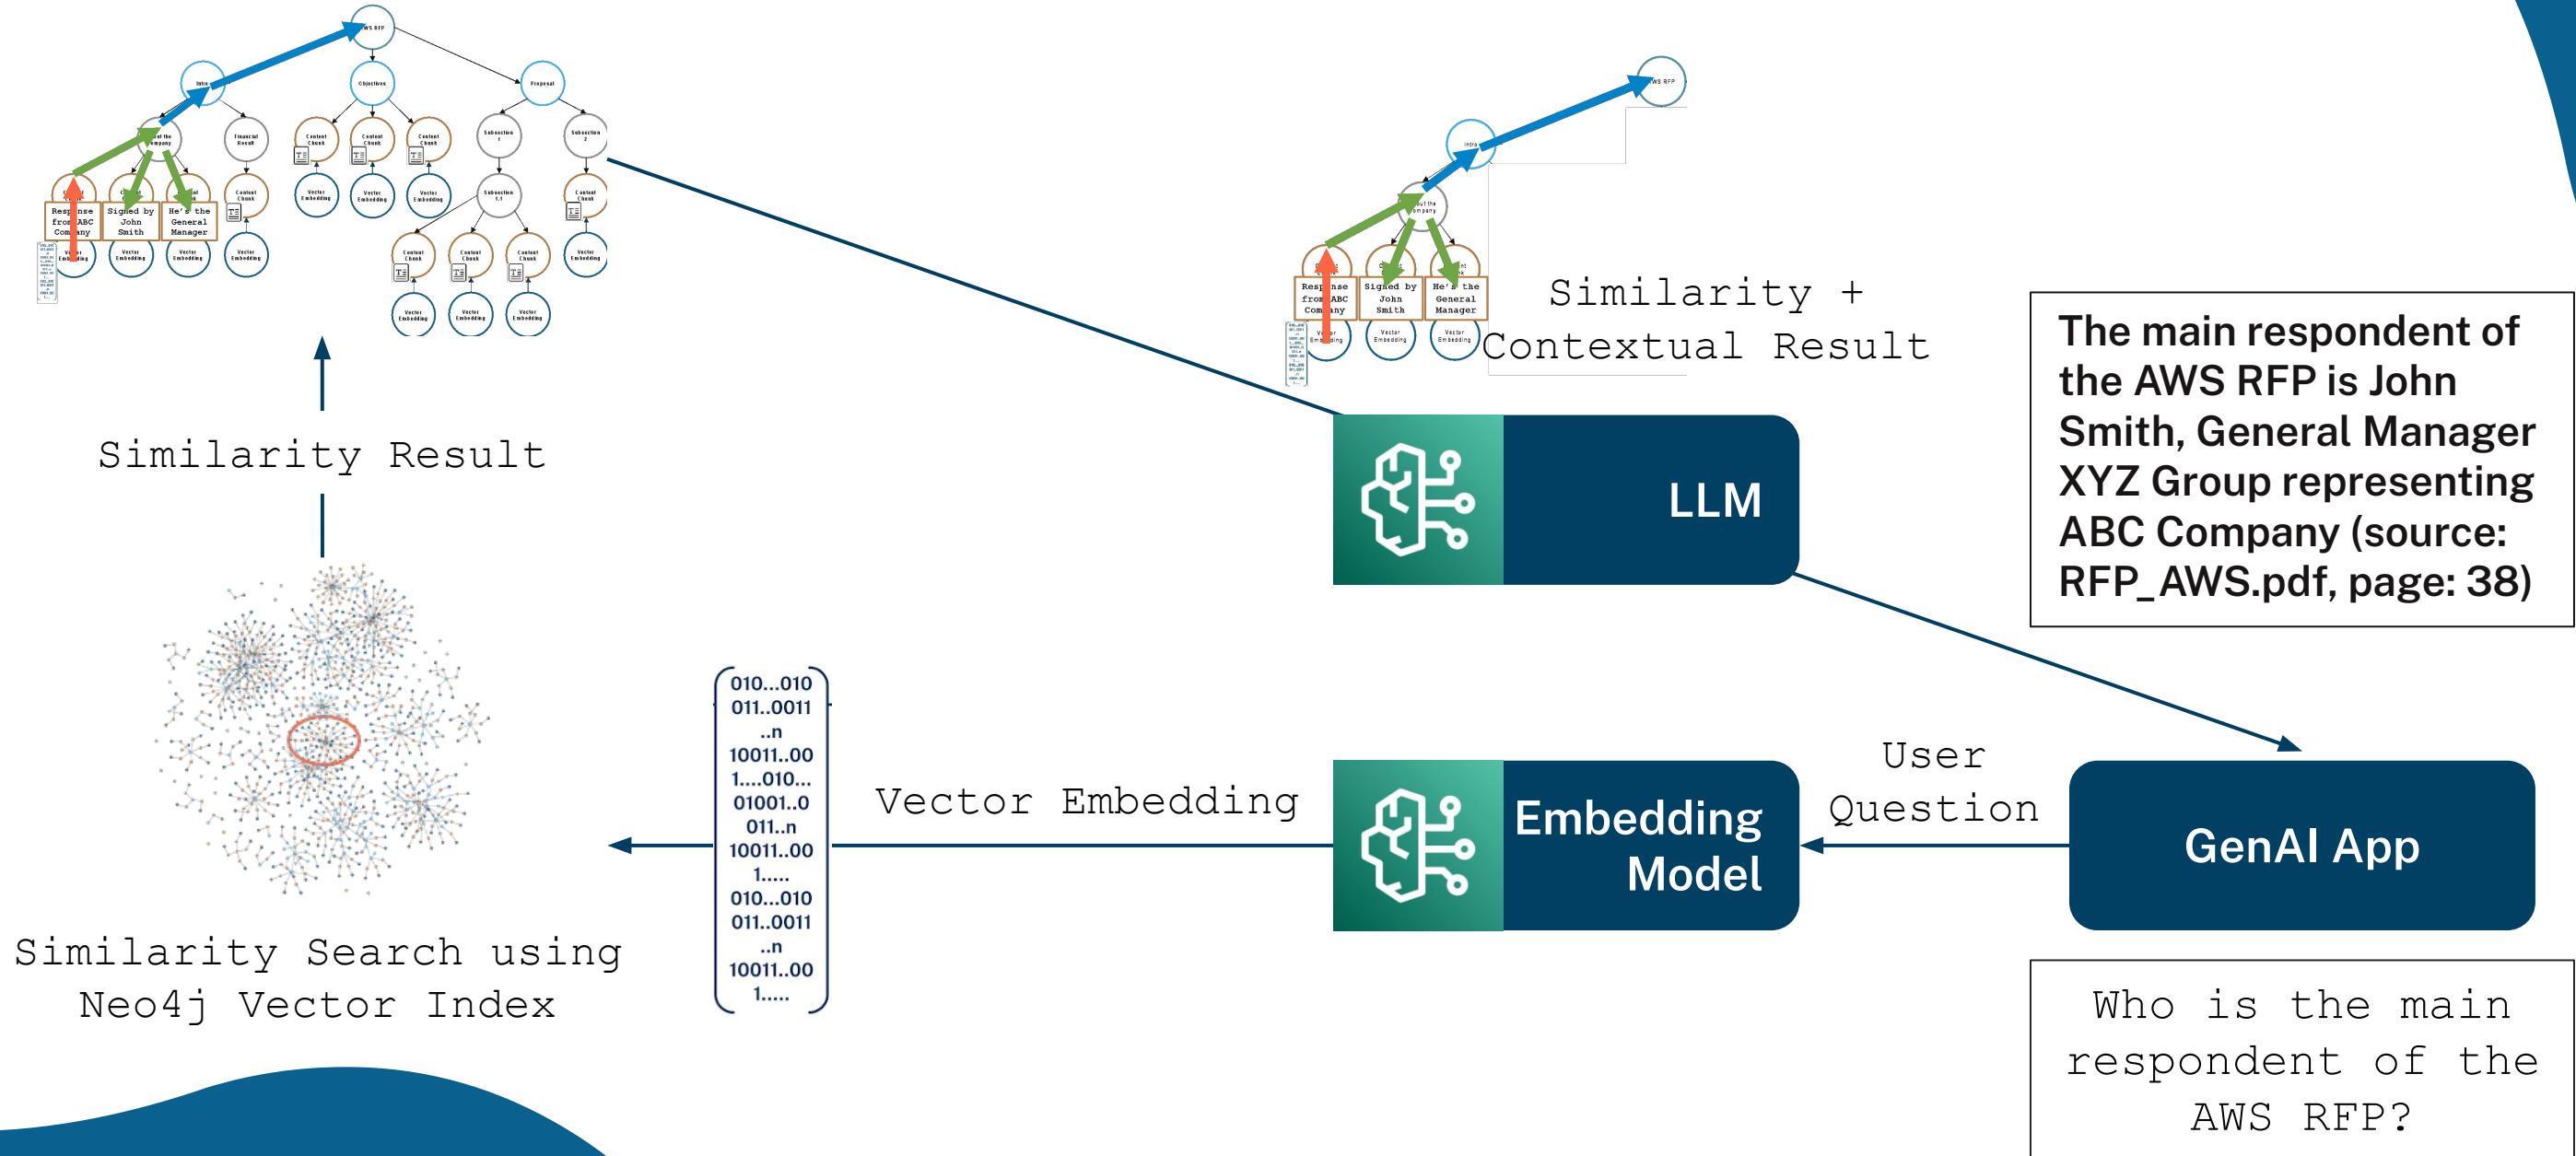
• Accurate, Contextual and Explainable



Accurate, Contextual and Explainable

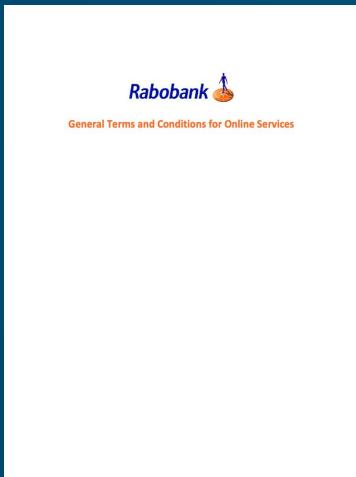


Accurate, Contextual and Explainable

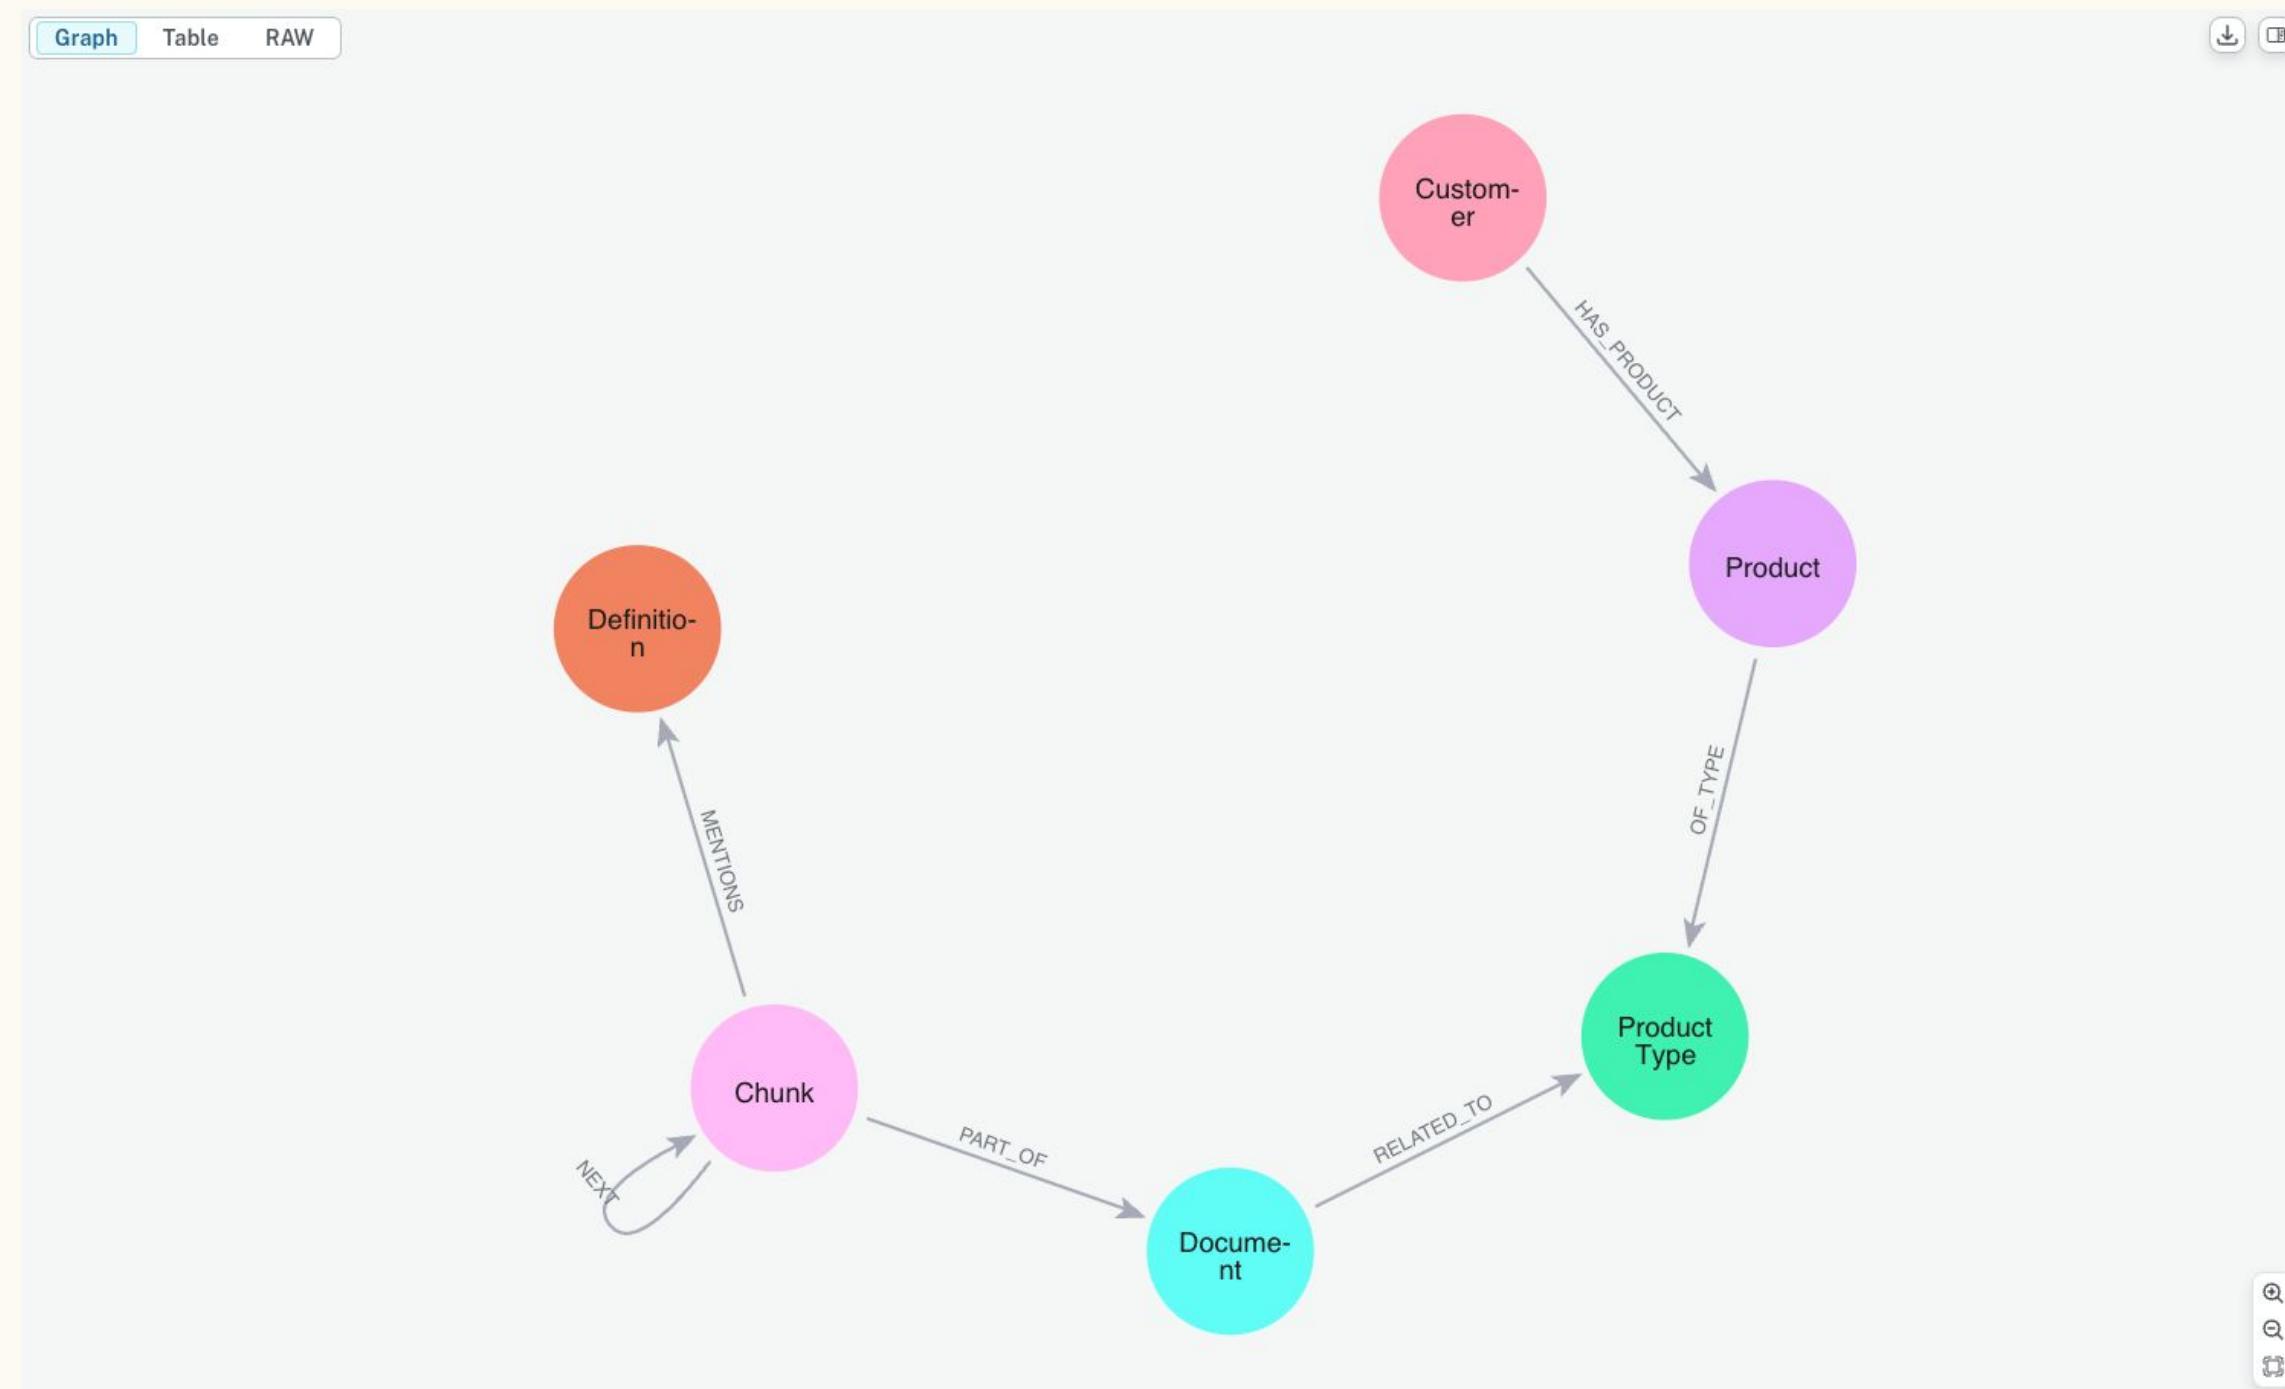


Rabobank Products

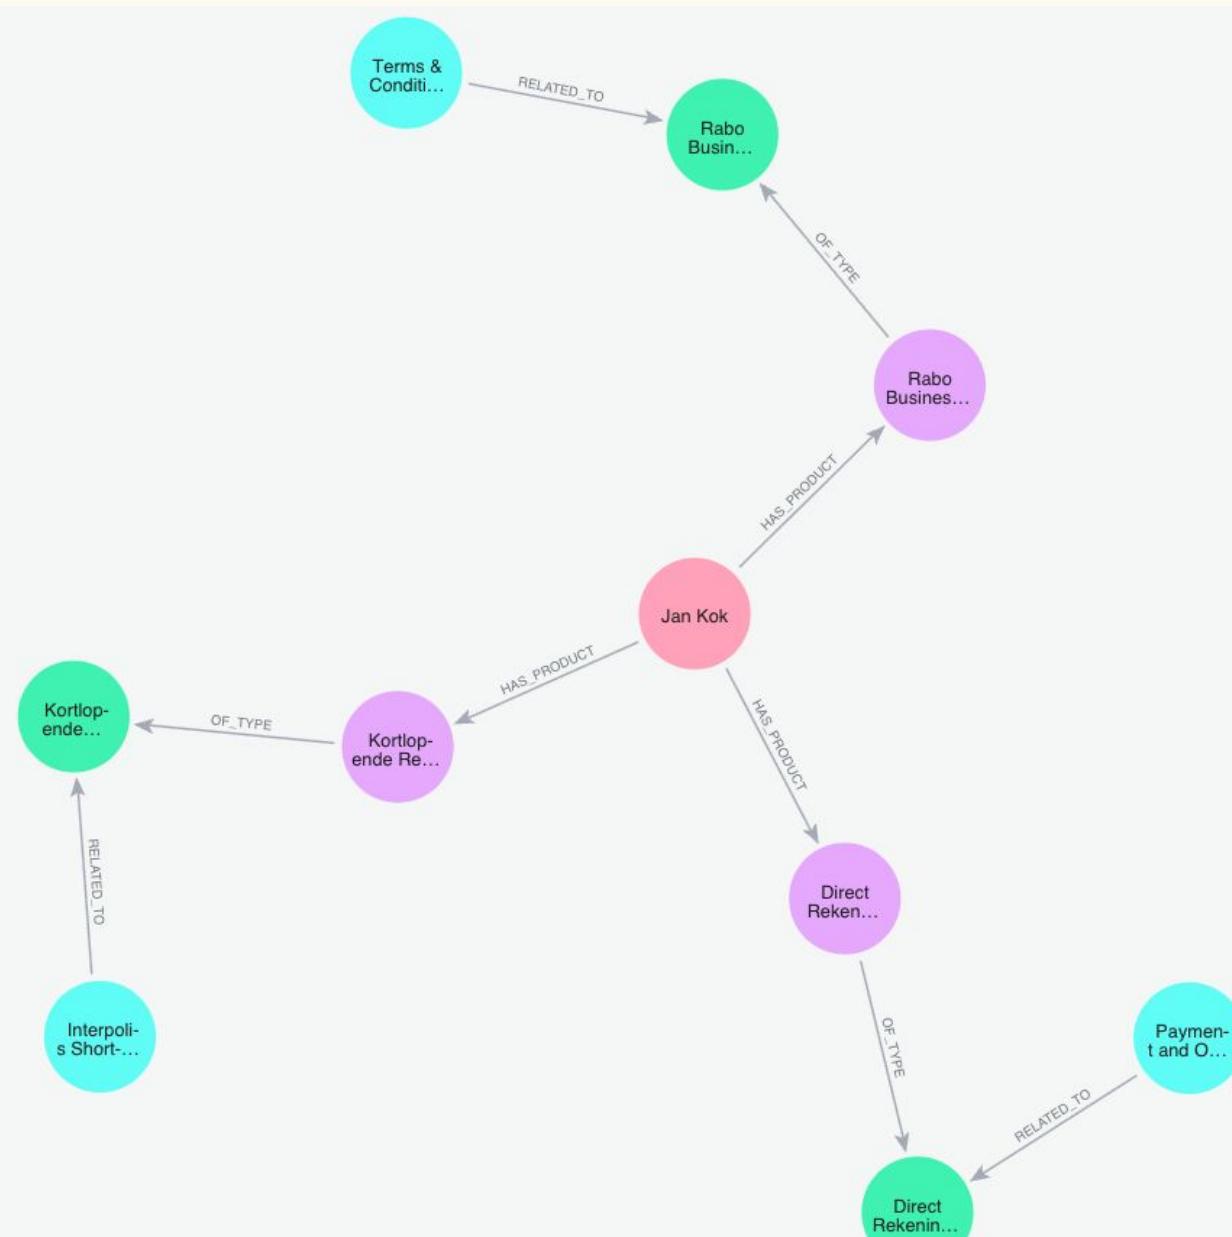
1. **Savings Account:** [Document from Rabo Spaarrekening.](#)
2. **Current Account:** [Document from Rabo DirectPakket.](#)
3. **Investment Account:** [Document from Effectenrekening en Rabo BeleggersRekening.](#)
4. **Travel Insurance:** [Document from Kortlopende Reisverzekering.](#)
5. **Online Business Services:** [Document from Rabo Business Banking.](#)



Our Data Model Today



Our Data Model Today



```
1 MATCH p=(n:Customer{name: "Jan Kok"})-[:HAS_PRODUCT]->(:Product)-[:OF_TYPE]->(:ProductType)<-[:RELATED_TO]-(:Document)
2 RETURN p
```

Module 1

Go to Notebook



Module 2

Vector Index: Set up the vector Search

Knowledge Graphs – New & Improved!

NOW WITH VECTORS!



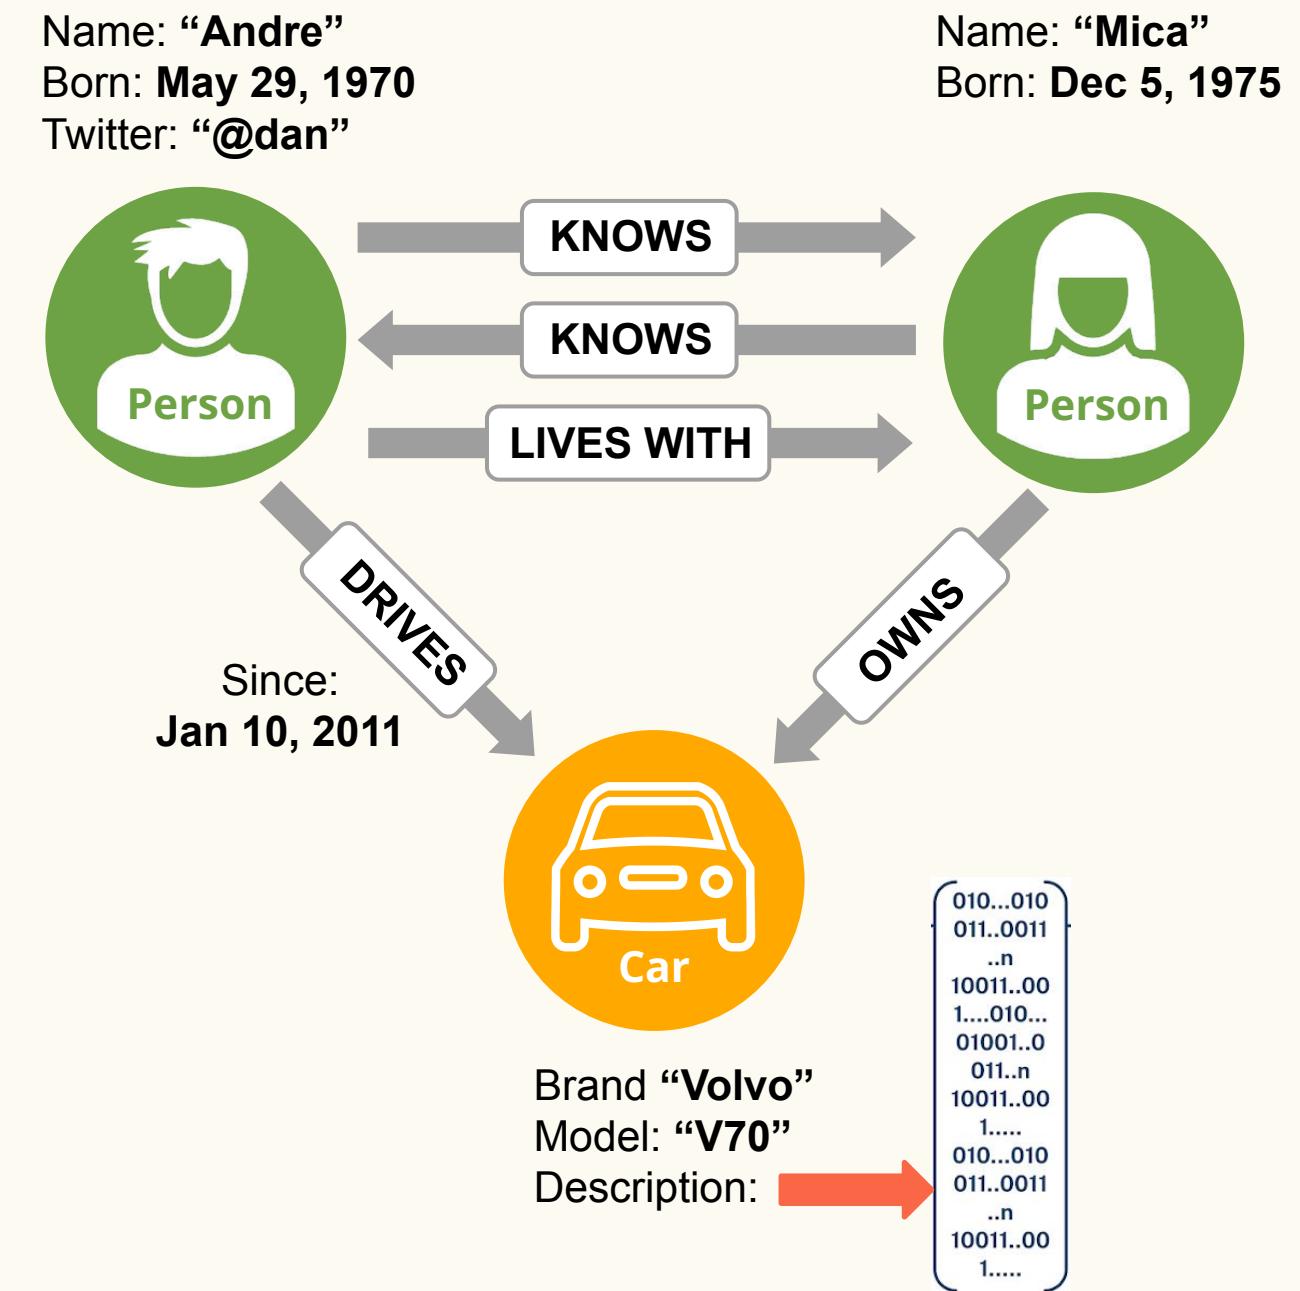
Knowledge Graph = design patterns to organize & access interrelated data

Property Graph Data Model

Nodes represent entities in the graph

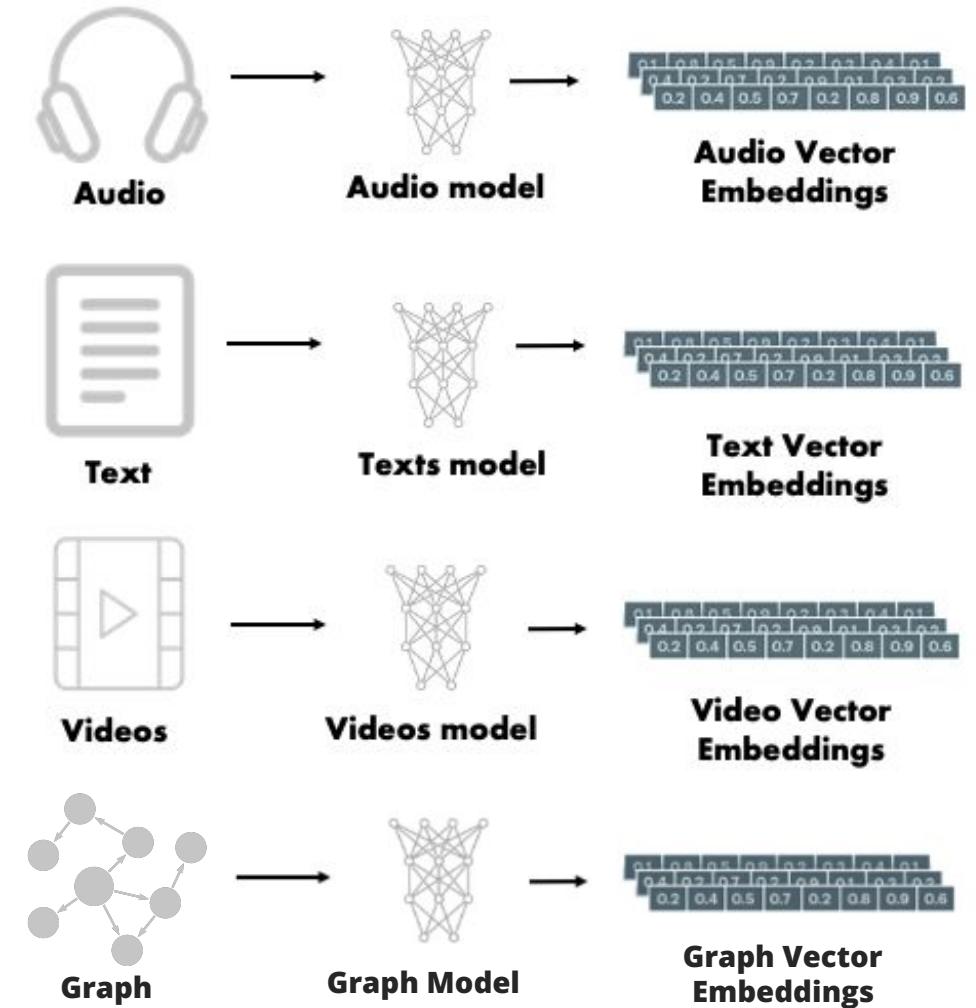
Relationships represent associations or interactions between nodes

Properties represent attributes of nodes or relationships



What are Embeddings?

- A type of data compression
- Transform messy data (text, images, audio, etc.) into a compact format for ML algorithms
- Most often numeric **vectors** (a.k.a arrays) with 100s or 1000s of elements
- Preserve information such that “*similar*” items *have proportionally “similar” embedding vectors*
- Similarity is measured with vector algorithms (cosine, euclidean, etc.)



Search & Vectors in Neo4j

Neo4j makes search efficient through a variety of indexes including

- **Range:** General index for predicates based on equality and range.
- **Point:** Predicates on geospatial points like distance bounding boxes, etc.
- **Text:** Predicates on strings like contains, ends with, etc.
- **FullText:** Text search based on tokenization and analyzers
- **Vector: ANN (Approximate Nearest Neighbor) search on vectors**

How is similarity measured

What would you use for openai embedding models?

For OpenAI embedding models, **cosine similarity** is the preferred choice. Here's why:

Why Cosine Similarity?

1. OpenAI Embeddings Are Normalized

- OpenAI embedding models (like `text-embedding-ada-002`) generate **normalized vectors**, meaning each vector has a unit length (i.e., a magnitude of 1).
- In such cases, **cosine similarity and Euclidean distance are proportional**, but cosine similarity is computationally more efficient.

2. Focus on Semantic Meaning

- Since cosine similarity measures the **angle between vectors**, it captures semantic similarity between words, sentences, or documents effectively.
- Euclidean distance, on the other hand, is sensitive to magnitude differences, which aren't meaningful in normalized embeddings.

3. Industry Best Practice

- Most applications using OpenAI embeddings (e.g., search ranking, recommendation systems, and semantic clustering) rely on cosine similarity.
- OpenAI's own documentation recommends using **cosine similarity or dot product over Euclidean distance**.

Module 3

Graph Analytics: Run some Graph Algorithms

50+ Graph Algorithms in Neo4j



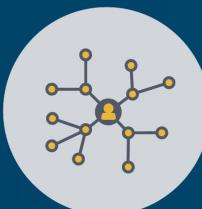
Pathfinding & Search

- Shortest Path
- Single-Source Shortest Path
- All Pairs Shortest Path
- A* Shortest Path
- Yen's K Shortest Path
- Minimum Weight Spanning Tree
- K-Spanning Tree (MST)
- Random Walk
- Breadth & Depth First Search



Link Prediction

- Adamic Adar
- Common Neighbors
- Preferential Attachment
- Resource Allocations
- Same Community
- Total Neighbors



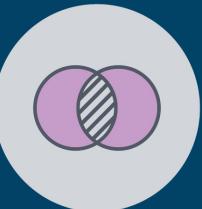
Centrality / Importance

- Degree Centrality
- Closeness Centrality
- Harmonic Centrality
- Betweenness Centrality & Approx.
- PageRank
- Personalized PageRank
- ArticleRank
- Eigenvector Centrality



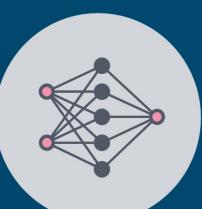
Community Detection

- Triangle Count
- Local Clustering Coefficient
- Connected Components (Union Find)
- Strongly Connected Components
- Label Propagation
- Louvain Modularity
- K-1 Coloring
- Modularity Optimization



Similarity

- Euclidean Distance
- Cosine Similarity
- Node Similarity (Jaccard)
- Overlap Similarity
- Pearson Similarity
- Approximate KNN



Embeddings

- Node2Vec
- Random Projections
- GraphSAGE

... Auxiliary Functions:

- Random graph generation
- Graph export
- One hot encoding
- Distributions & metrics

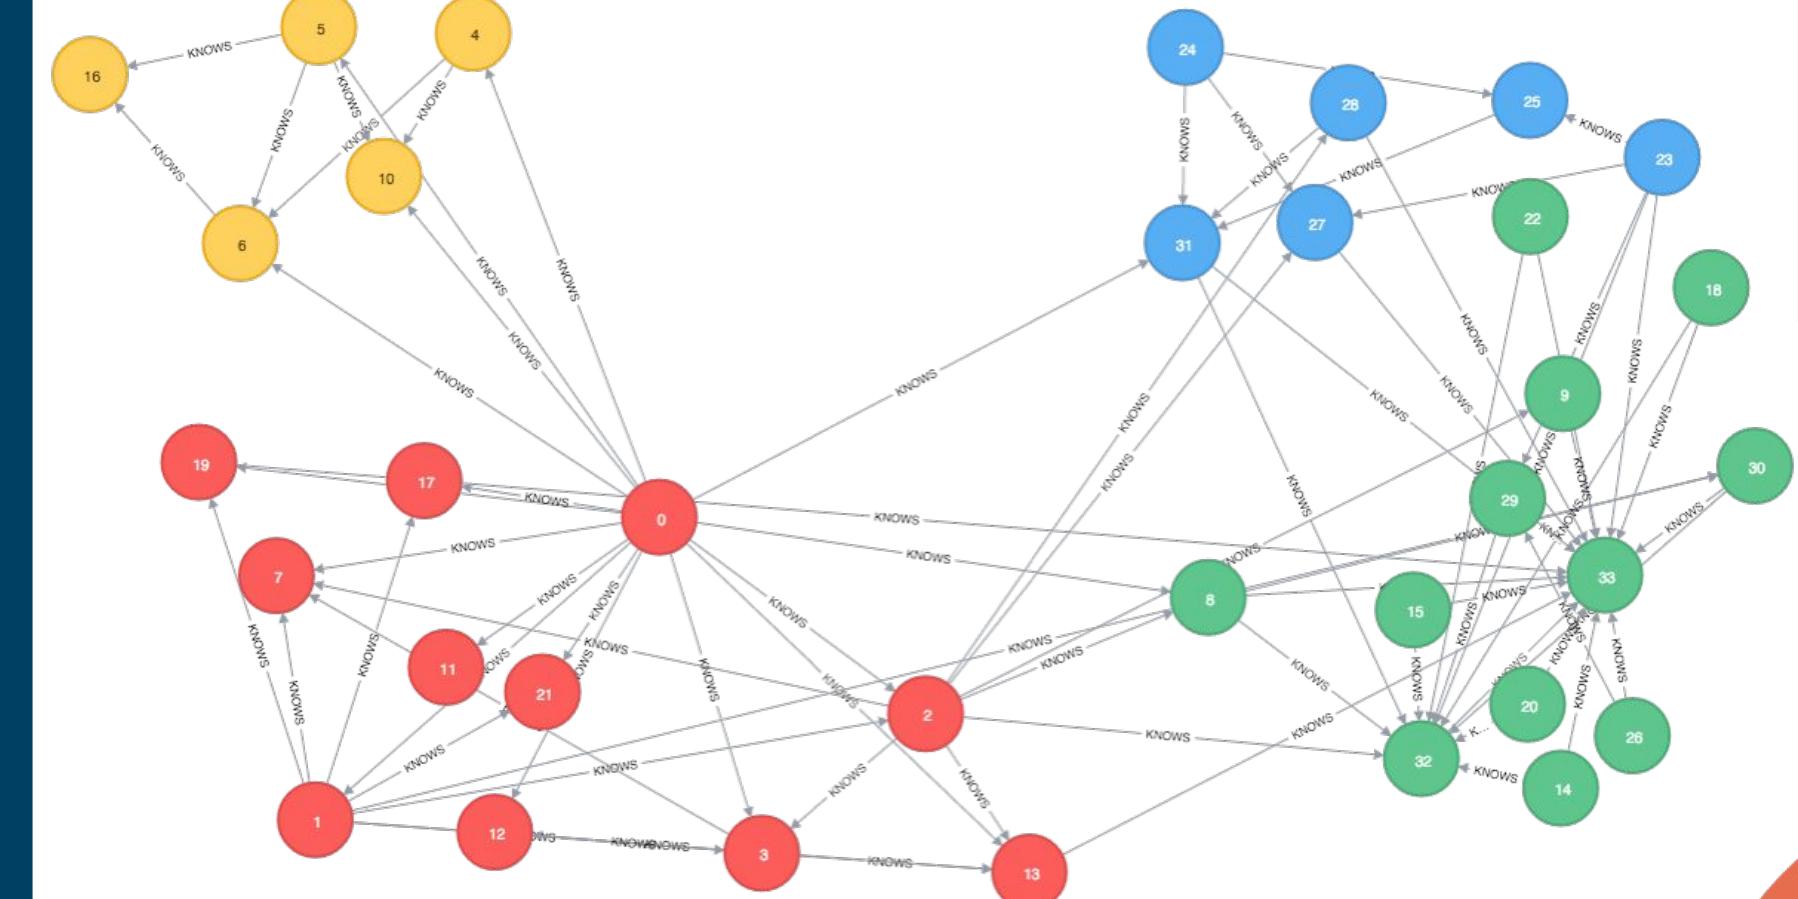
Today we will explore **Community Detection**

Evaluate how groups of nodes may be clustered or partitioned in a graph.

Community id properties assigned to node based on relationship structure.

Useful for:

- Segmentation
 - Clustering
 - Entity resolution
 - Summarization (for AI)



Module 4

GraphRAG Chatbot: Run a Chatbot on the Graph

Knowledge Graphs + LLM: Bringing it Together



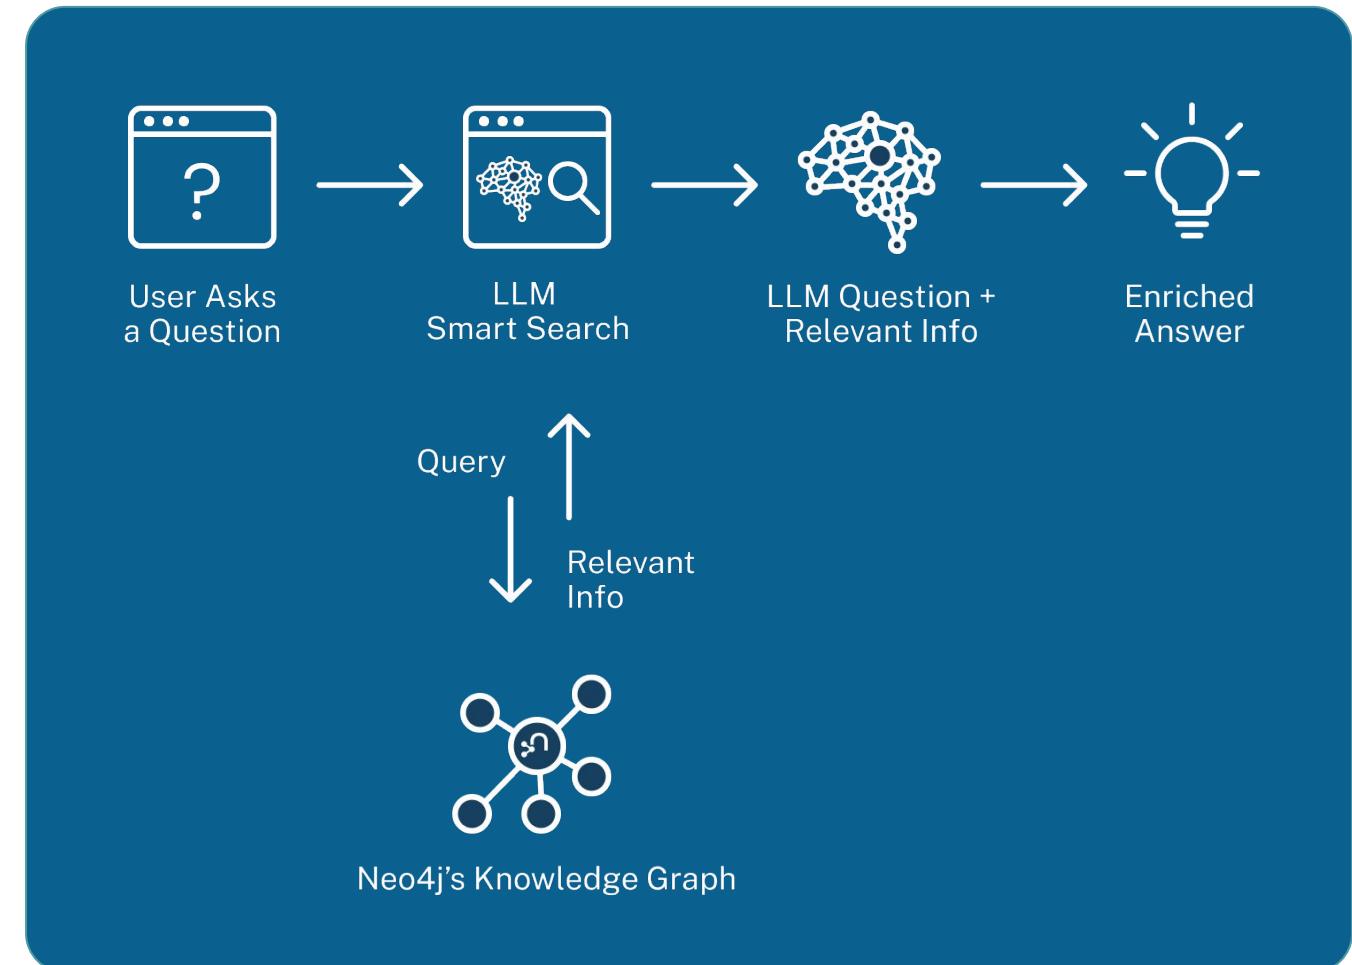
Building a RAG-application in Neo4j

Preparation Steps:

1. Chunk Documents
2. Create Embeddings
3. Load Chunks & Embeddings to the Database
4. Create a Vector Index

RAG Steps:

1. User Provides the User Query
2. Embed the User Query
3. Retrieve Relevant Documents
4. Prompt the LLM with User Query and Relevant Documents
5. Provide Answer to User



GraphRAG

From Local to Global: A Graph RAG Approach to Query-Focused Summarization

Darren Edge^{1†} Ha Trinh^{1†} Newman Cheng² Joshua Bradley² Alex Chao³

Apurva Mody³

Steven Truitt²

Jonathan Larson¹

¹Microsoft Research

²Microsoft Strategic Missions and Technologies

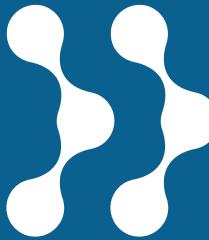
³Microsoft Office of the CTO

{daedge, trinhha, newmarcheng, joshbradley, achao, moapurva, steventruitt, jolarso}
@microsoft.com

[†]These authors contributed equally to this work

Abstract

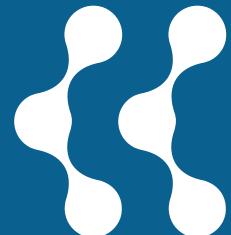
The use of retrieval-augmented generation (RAG) to retrieve relevant information from an external knowledge source enables large language models (LLMs) to answer questions over private and/or previously unseen document collections. However, RAG fails on global questions directed at an entire text corpus, such

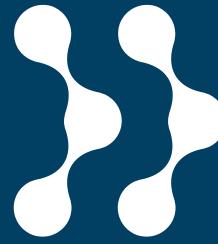


GraphRAG

*Advanced RAG Patterns that
use Graph Data Structures for
Retrieval for relevant context
and higher explainability.*

Patterns





arXiv:2404.17723v2 [cs.IR] 6 May 2024

Retrieval-Augmented Generation with Knowledge Graphs for Customer Service Question Answering

Zhentao Xu
zhxw@linkedin.com
LinkedIn Corporation
Sunnyvale, CA, USA

Tie Wang
tiewang@linkedin.com
LinkedIn Corporation
Sunnyvale, CA, USA

Mark Jerome Cruz
marcruz@linkedin.com
LinkedIn Corporation
Sunnyvale, CA, USA

Manasi Deshpande
madeshpande@linkedin.com
LinkedIn Corporation
Sunnyvale, CA, USA

Matthew Guevara
mguevara@linkedin.com
LinkedIn Corporation
Sunnyvale, CA, USA

Xiaofeng Wang
xiaofwang@linkedin.com
LinkedIn Corporation
Sunnyvale, CA, USA

Zheng Li
zelij@linkedin.com
LinkedIn Corporation
Sunnyvale, CA, USA

KEYWORDS
Large Language Model, Knowledge Graph, Question Answering, Retrieval-Augmented Generation

ACM Reference Format:
Zhentao Xu, Mark Jerome Cruz, Matthew Guevara, Tie Wang, Manasi Deshpande, Xiaofeng Wang, and Zheng Li. 2024. Retrieval-Augmented Generation with Knowledge Graphs for Customer Service Question Answering. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval (SIGIR '24)*, July 14–18, 2024, Washington, DC, USA. ACM, New York, NY, USA, 5 pages. <https://doi.org/10.1145/3626772.3661370>

1 INTRODUCTION
Effective technical support in customer service underpins product success, directly influencing customer satisfaction and loyalty. Given the frequent similarity of customer inquiries to previously resolved issues, the rapid and accurate retrieval of relevant past instances is crucial for the efficient resolution of such inquiries. Recent advancements in embedding-based retrieval (EBR), large language models (LLMs), and retrieval-augmented generation (RAG) [8] have significantly enhanced retrieval performance and question-answering capabilities for the technical support of customer service. This process typically unfolds in two stages: first, historical issue tickets are treated as plain text, segmented into smaller chunks to accommodate the context length constraints of embedding models; each chunk is then converted into an embedding vector for retrieval. Second, during the question-answering phase, the system retrieves the most relevant chunks and feeds them as contexts for LLMs to generate answers in response to queries. Despite its straightforward approach, this method encounters several limitations:

- **Limitation 1 - Compromised Retrieval Accuracy from Ignoring Structures:** Issue tracking documents such as Jira [2] possess inherent structure and are interconnected, with references such as "issue A is related to/copied from/caused

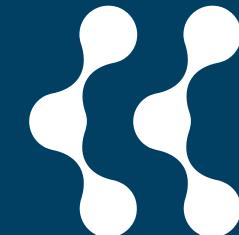
Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyright for works published outside the US may be held by the author or the author's institution. Advertising with the publisher for profit is prohibited. Requests for permission should be addressed to permissions@acm.org.

SIGIR '24, July 14–18, 2024, Washington, DC, USA
© 2024 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ISBN 978-1-4503-6613-7
<https://doi.org/10.1145/3626772.3661370>

We introduce a novel customer service question-answering method that **amalgamates RAG with a knowledge graph (KG)**. (...)

Empirical assessments on our benchmark datasets, utilizing key retrieval and text generation metrics, **reveal that our method outperforms the baseline by 77.6% in MRR and by 0.32 in BLEU**. Our method has been deployed within LinkedIn's customer service team for approximately six months and has reduced the median per-issue resolution time by 28.6%.

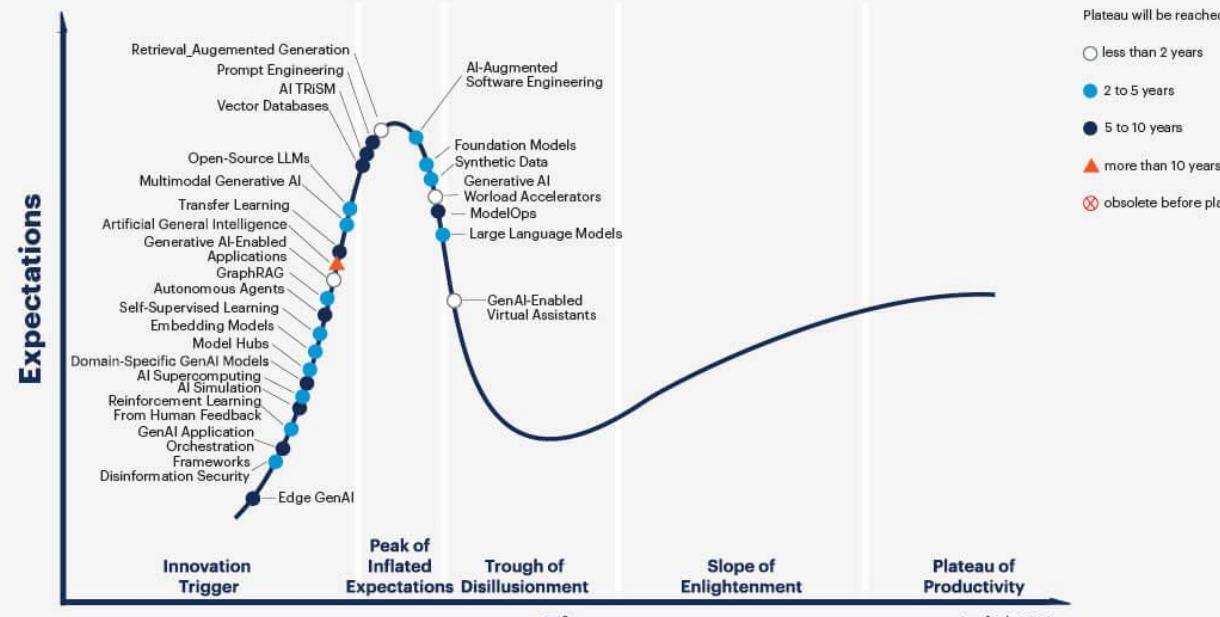
Retrieval-Augmented Generation with Knowledge Graphs for Customer Service Question Answering -LinkedIn





Gartner®

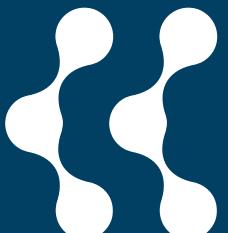
Hype Cycle for Generative AI, 2024



GraphRAG: This technique improves the accuracy, reliability and explainability of retrieval-augmented generation (RAG) systems. The approach uses knowledge graphs (KGs) to improve the recall and precision of retrieval, either directly by pulling facts from a KG or indirectly by optimizing other retrieval methods.

Gartner®

Gartner Hype Cycle Generative AI, November 2024



Example Pattern

Name: Graph Enhanced Vector Search

Description: The user question is embedded using the same embedder used to create chunk embeddings. A vector similarity search is executed on the chunk embeddings to find k (number previously configured by developer/user) most similar chunks. A traversal of the Domain Graph starting at the found chunks is executed to retrieve more context.

Context: The biggest problem with basic GraphRAG patterns is finding all relevant context necessary to answer a question. The context can be spread across many chunks not being found by the search. Relating the real-world entities from the chunks to each other and retrieving these relationships together with a vector search provides additional context about these entities that the chunks refer to. They can also be used to relate chunks to each other through the entity network.

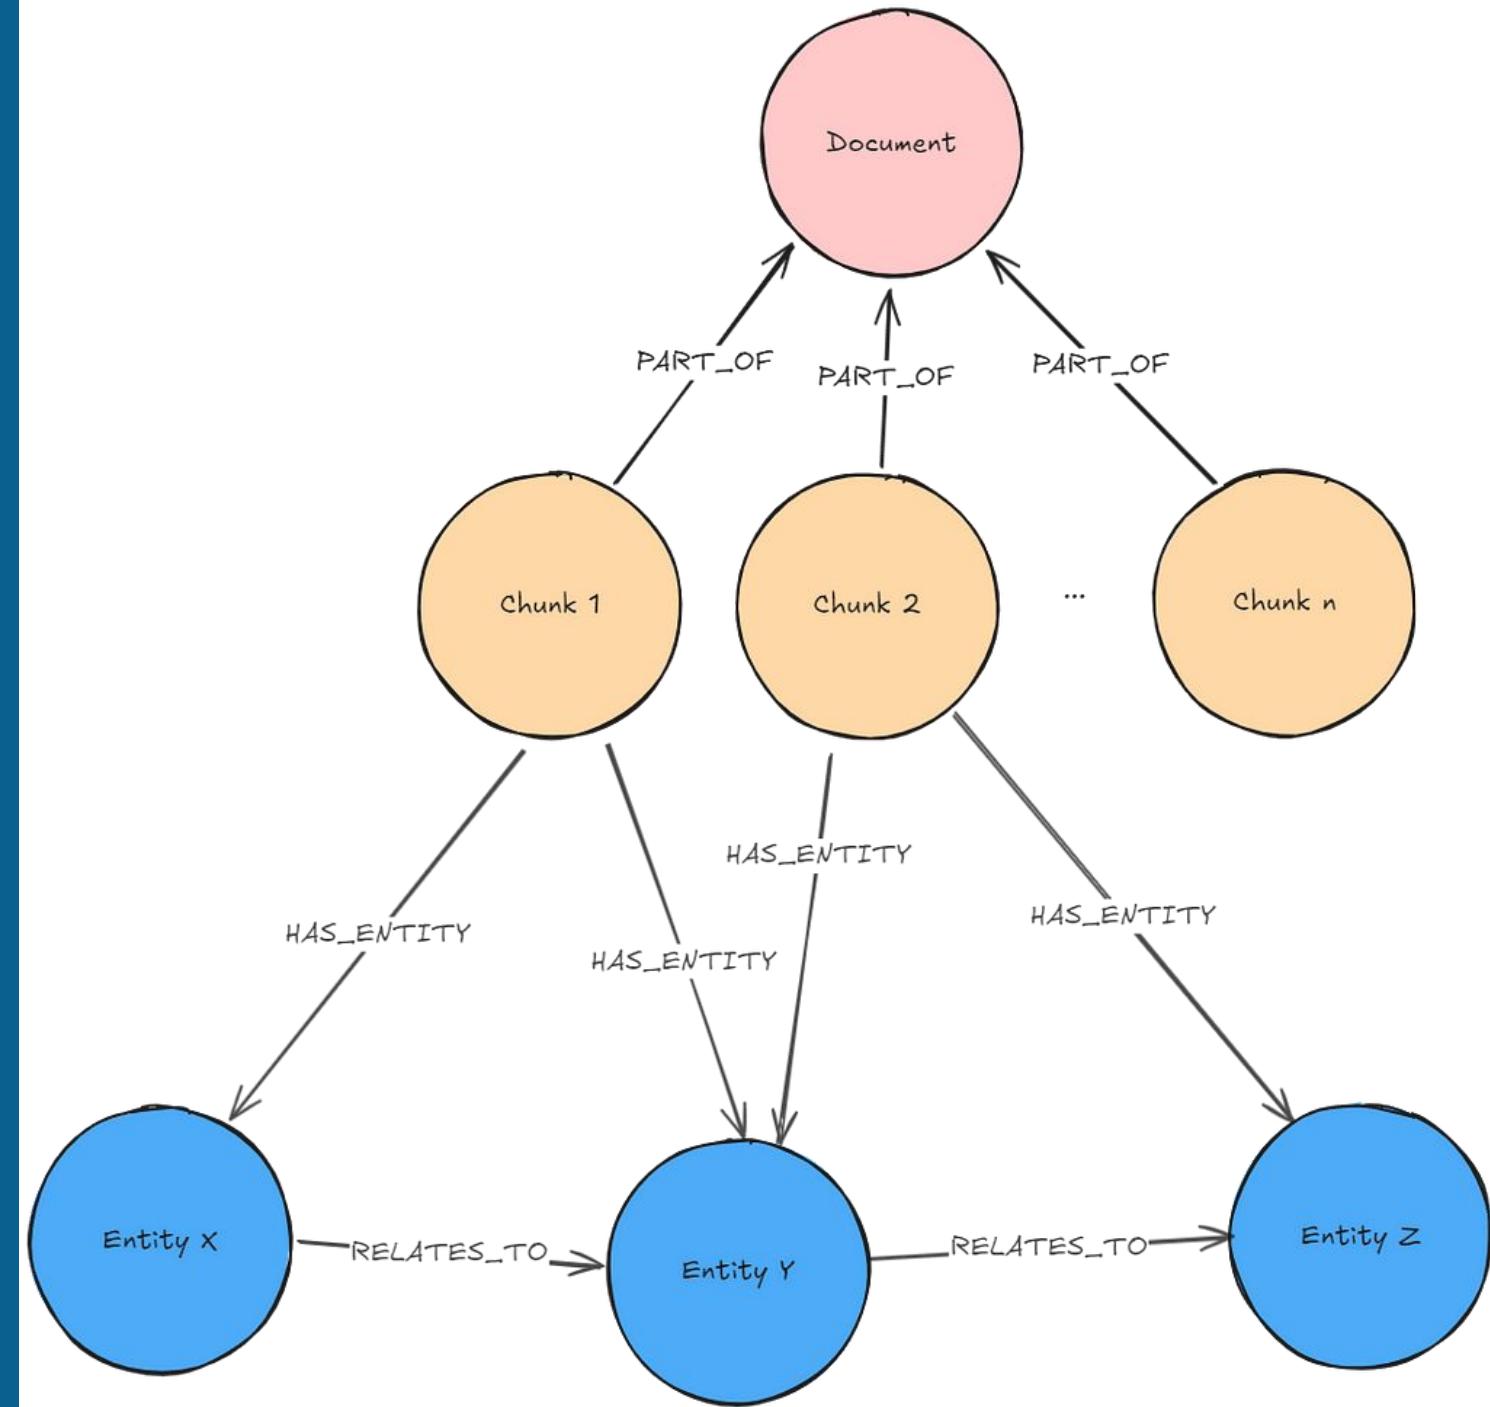
Required pre-processing: Use an LLM to execute entity and relationship extraction on the chunks. Import the retrieved triples into the graph.

Variations: Entity disambiguation, Question-guided/Schema-defined extraction, Entity embeddings, Ontology-driven traversal

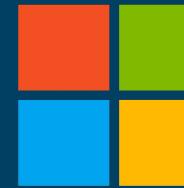
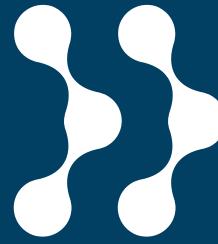
```
MATCH (node)-[:PART_OF]->(d:Document)
MATCH (node)-[:HAS_ENTITY]->(e)
MATCH path=(e)((()-[rels:!HAS_ENTITY&!PART_OF]-()){0,2}(:!Chunk&!Document))
...
RETURN ...
```

AKA: Graph + Vector, Augmented Vector Search

Required graph pattern: Lexical Graph with Extracted Entities



Graph Enhanced Vector Search



Microsoft

GraphRAG: Unlocking LLM discovery on narrative private data

Published February 13, 2024

By Jonathan Larson, Senior Principal Data Architect; Steven Truitt, Principal Program Manager

Share this page [f](#) [t](#) [l](#) [g](#) [n](#)



Perhaps the greatest challenge – and opportunity – of LLMs is extending their powerful capabilities to solve problems beyond the data on which they have been trained, and to achieve comparable results with data the LLM has never seen. This opens new possibilities in data investigation, such as identifying themes and semantic concepts with context and grounding on datasets. In this post, we introduce GraphRAG, created by Microsoft Research, as a significant advance in enhancing the capability of LLMs.

Retrieval-Augmented Generation (RAG) is a technique to search for information based on a user query and provide the results as reference for an AI answer to be generated. This technique is an important part of most LLM-based tools and the majority of RAG approaches use vector similarity as the search technique. GraphRAG uses LLM-generated knowledge graphs to provide substantial improvements in question-and-answer performance when conducting document analysis of complex information. This builds upon our recent [research](#), which points to the power of prompt augmentation when performing discovery on *private datasets*. Here, we define *private dataset* as data that the LLM is not trained on and has never seen before, such as an enterprise's proprietary research, business documents, or communications. *Baseline RAG* was created to help solve this problem, but we observe situations where baseline RAG performs very poorly. For example:

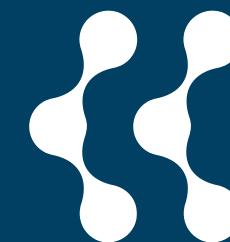
- Baseline RAG struggles to connect the dots. This happens when answering a question requires traversing disparate pieces of information through their shared attributes in order to provide new synthesized insights.
- Baseline RAG performs poorly when being asked to holistically understand summarized semantic concepts over large data collections or even singular large documents.

*Initial results show that GraphRAG **consistently outperforms** baseline RAG.*

Steps:

1. *Ingest Text Data*
2. *Generate Knowledge Graph*
3. *Import Into Graph Database*
4. *Create Semantic Hierarchies*
5. *Augment Retrieval*
6. *Deeper Understanding*

[Microsoft - GraphRAG](#)



Global Community Summary Retriever

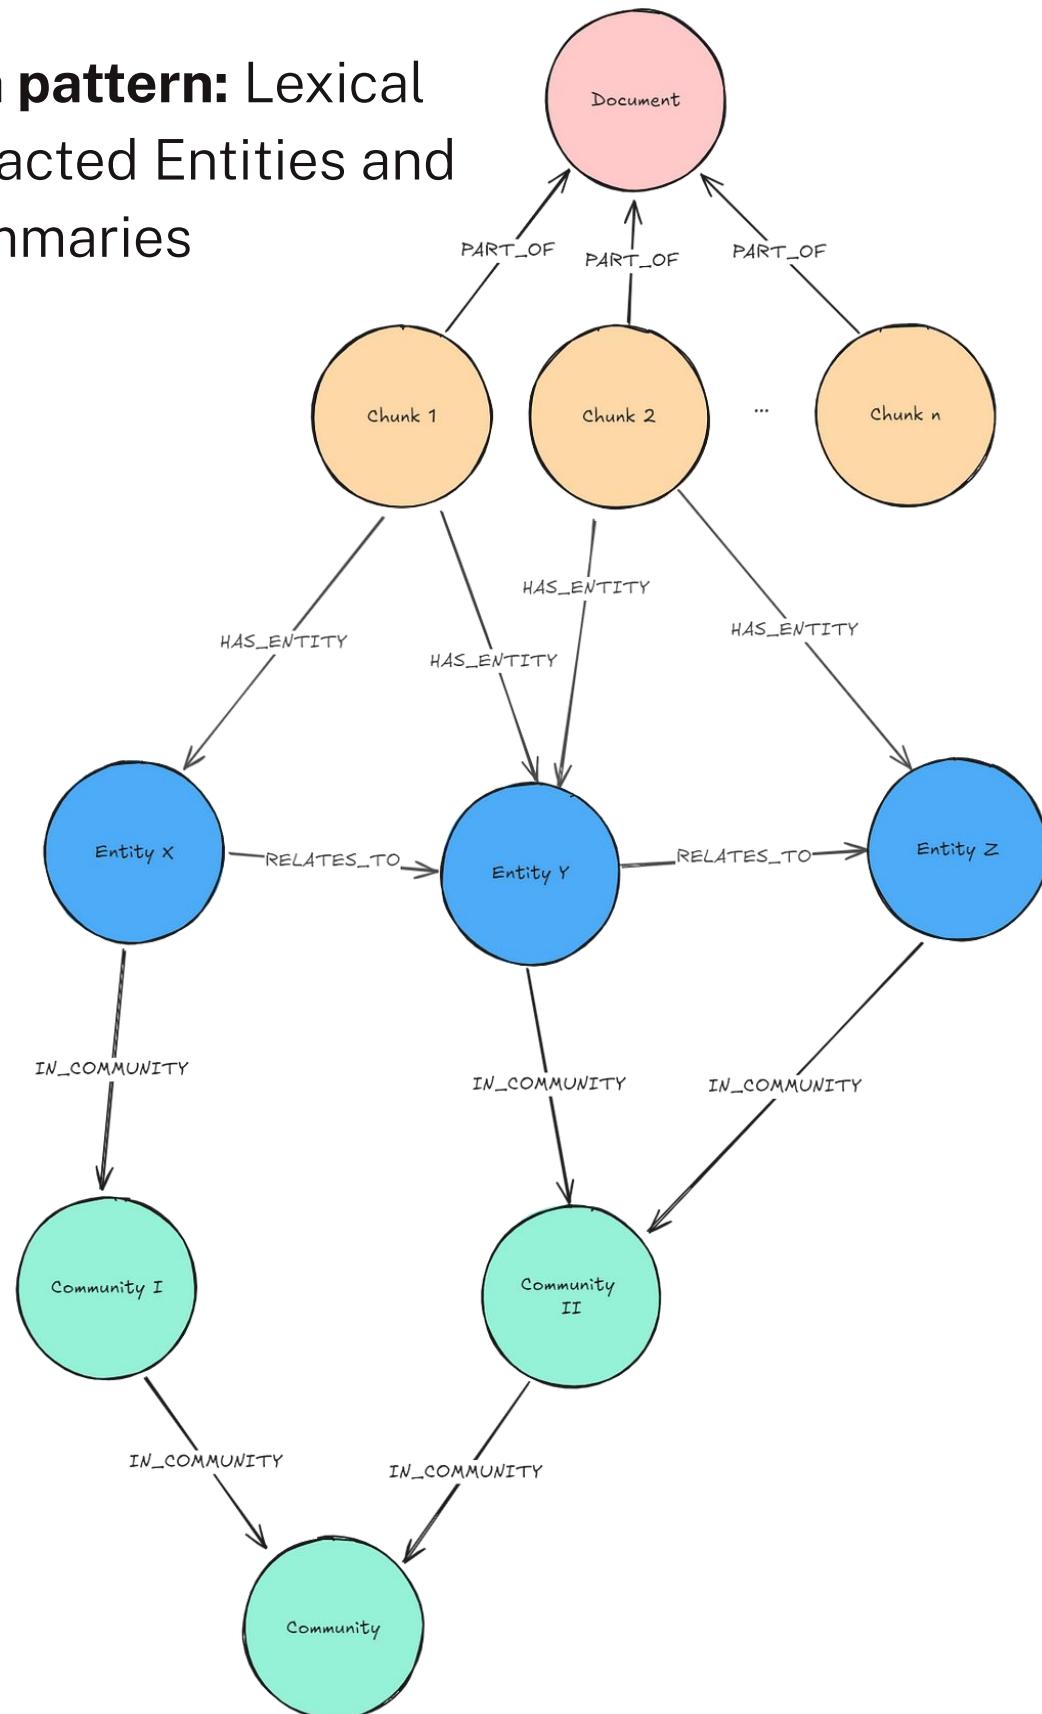
AKA: *Microsoft GraphRAG*, Global Retriever

Context: Certain *questions that can be asked on a whole dataset* do not just relate to things present in some chunks but rather search for an overall message that is overarching in the dataset.

Required pre-processing: In addition to extracting entities and their relationships, we need to form hierarchical communities within the Domain Graph. For every community, an LLM summarizes the entity and relationship information into Community Summaries.

```
MATCH (c:_Community_)  
WHERE c.level = $level  
RETURN c.full_content AS output
```

Required graph pattern: Lexical Graph with Extracted Entities and Community Summaries

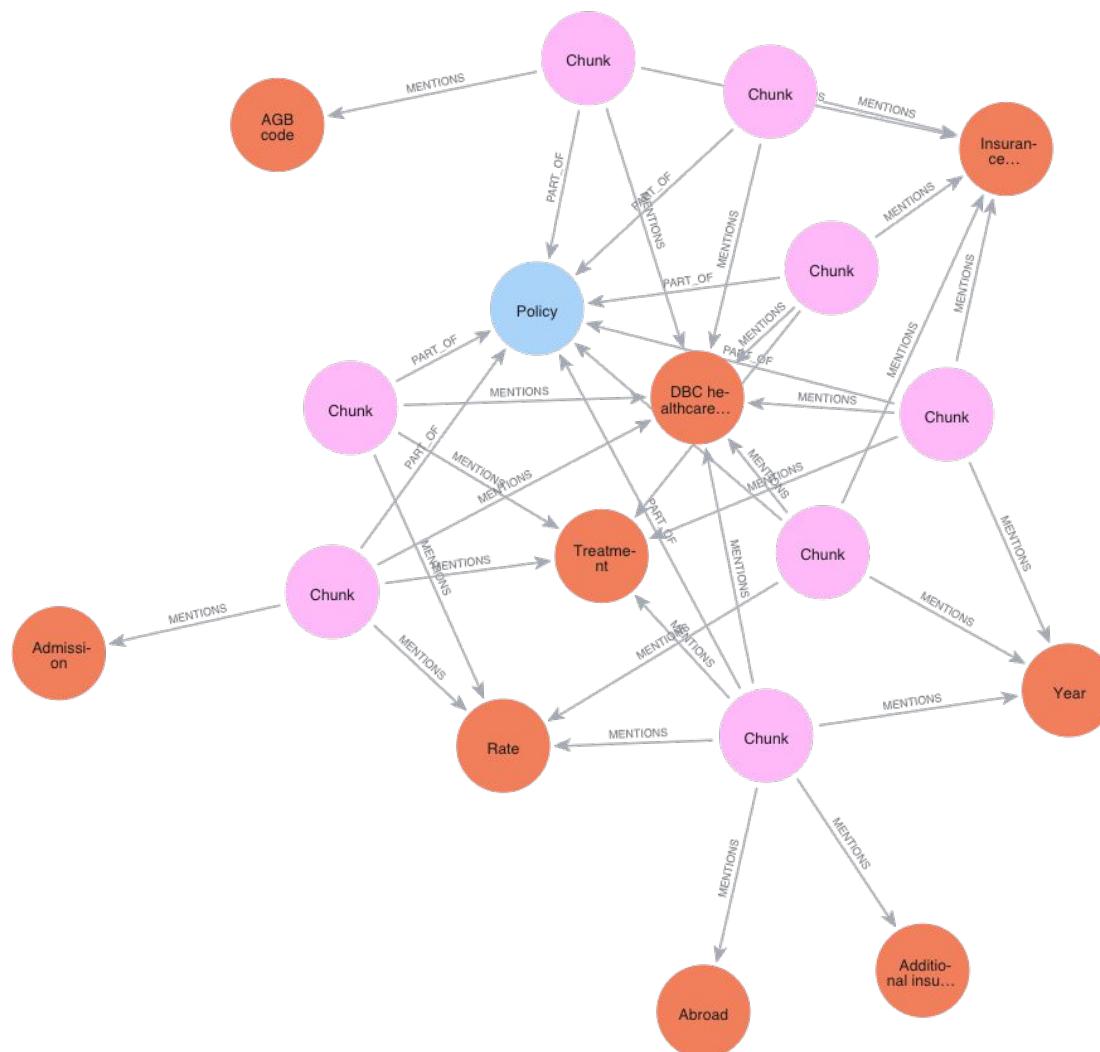


Add Domain Knowledge

- There are references in the document that give more context
- These add more context to the chunks



```
1 MATCH (c:Chunk)-[:MENTIONS]-(d:Definition)
2 WHERE c.id in $chunk_ids
3 WITH DISTINCT d as d
4 RETURN d.definition as definition, d.description as description
```



Module 4

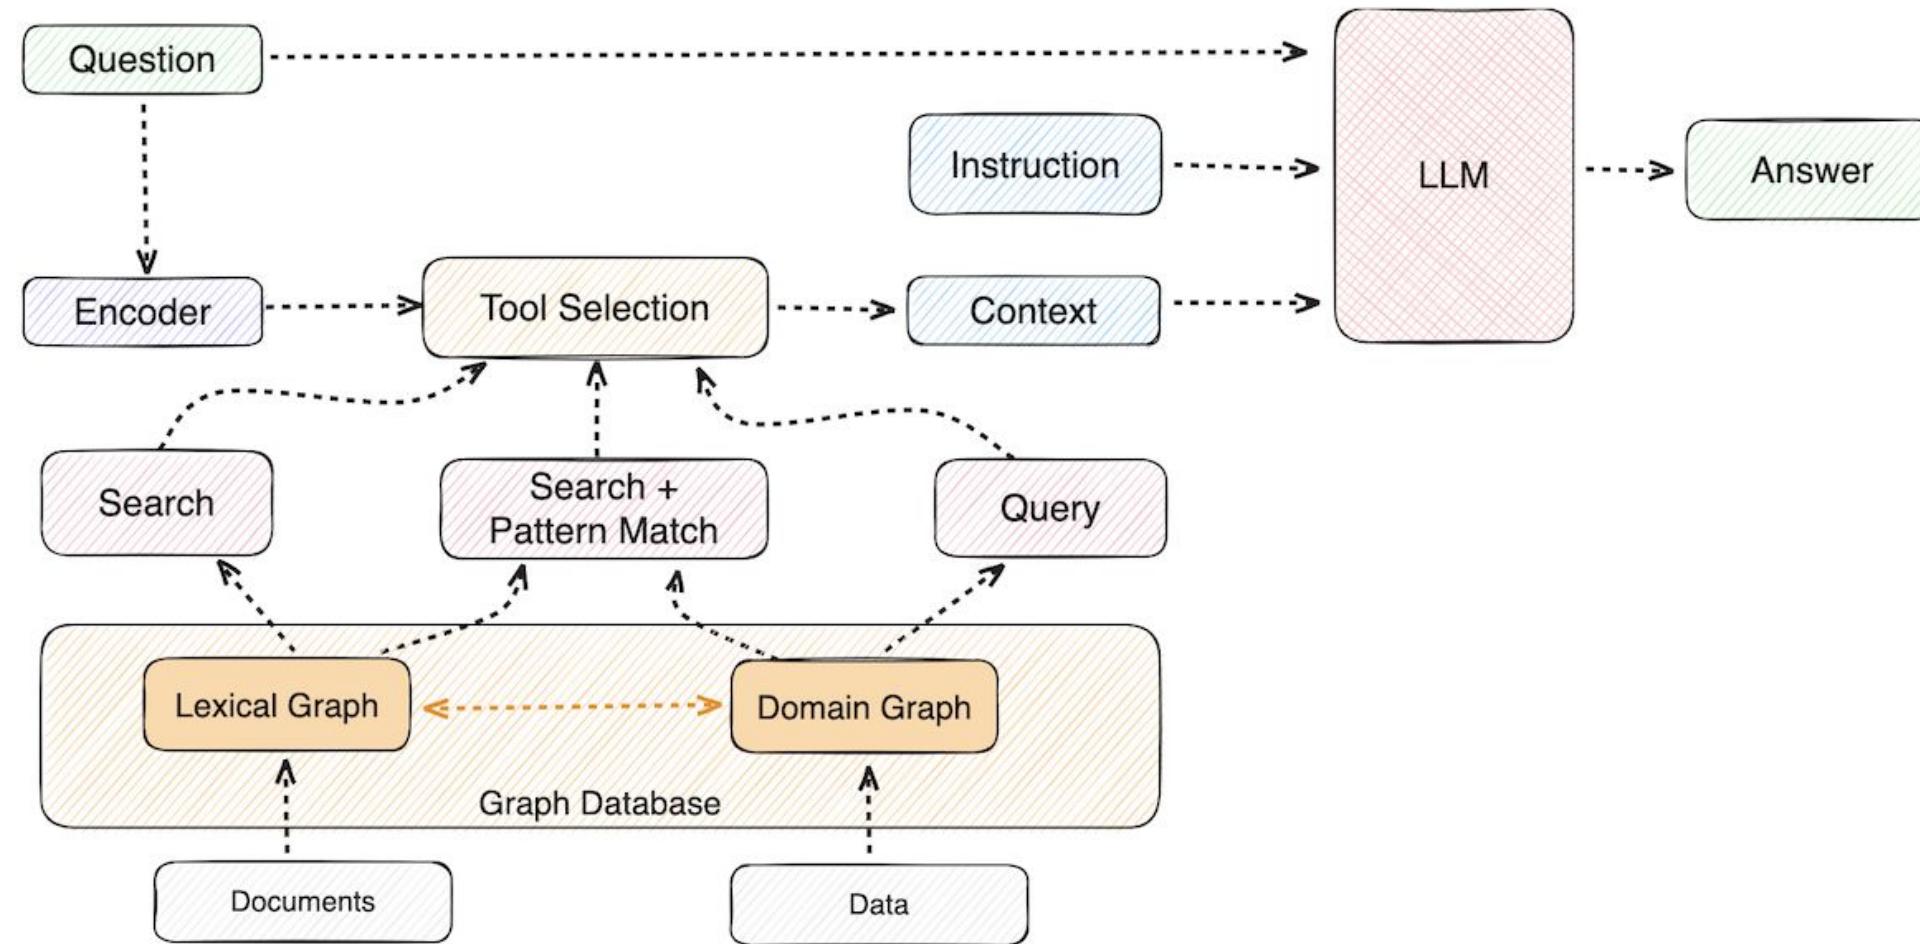
Go to Notebook



Module 5

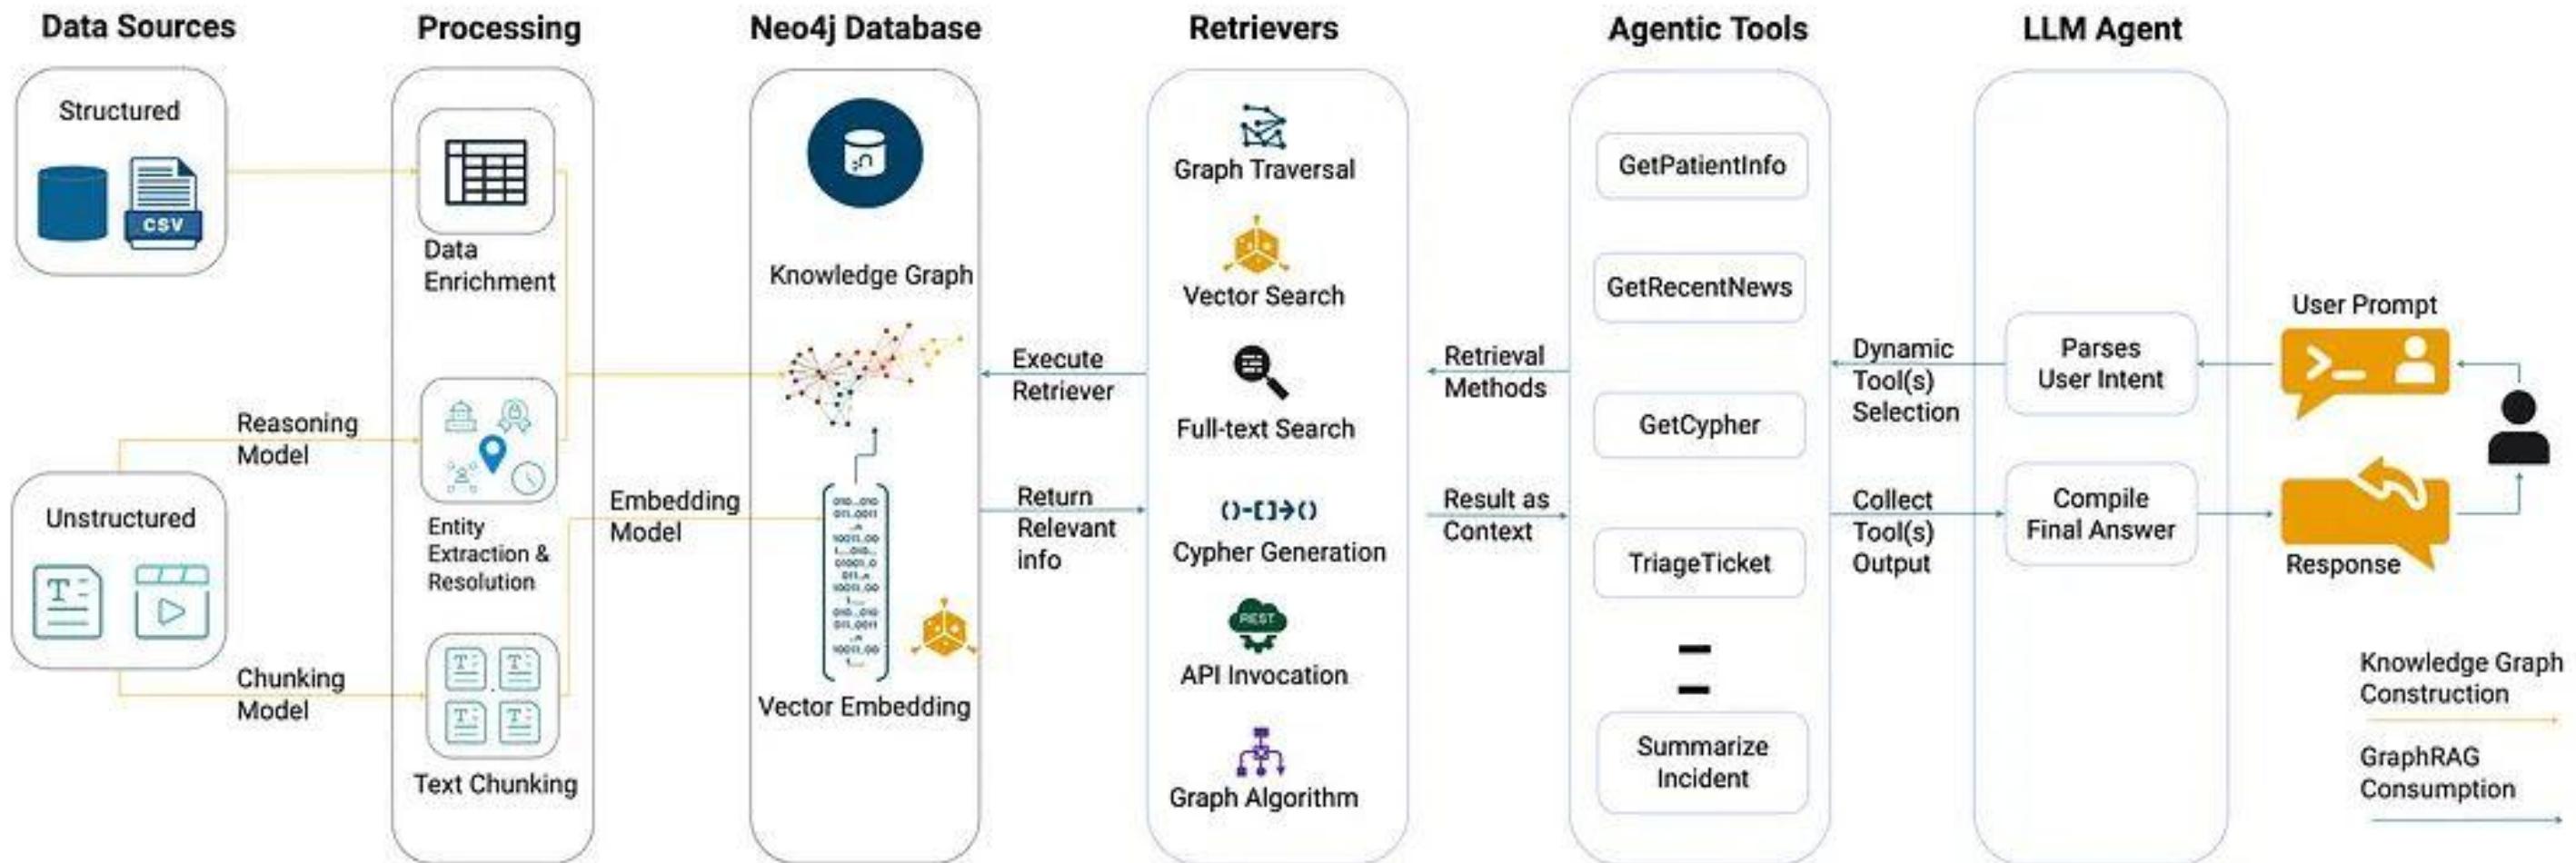
GraphRAG Agent: Create Agents with Tools

What is GraphRAG?



[Blog - What is GraphRAG?](#)

Agentic GraphRAG Architecture



Module 5

Go to Notebook

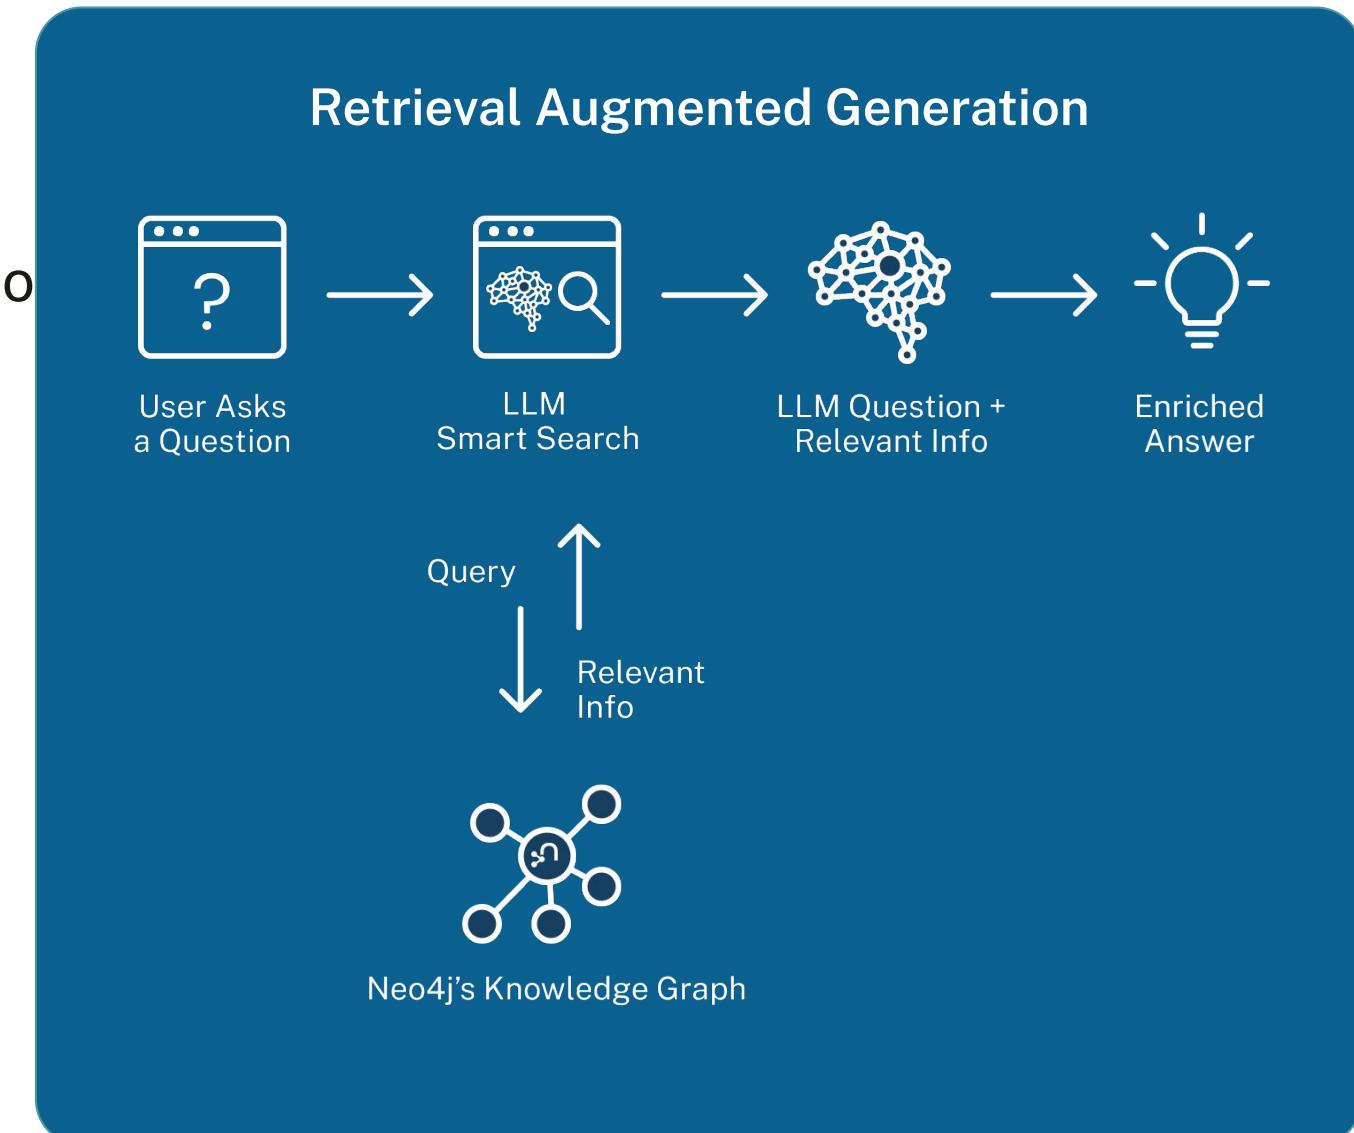


Wrap Up!

Neo4j Knowledge Graph as Database of Truth

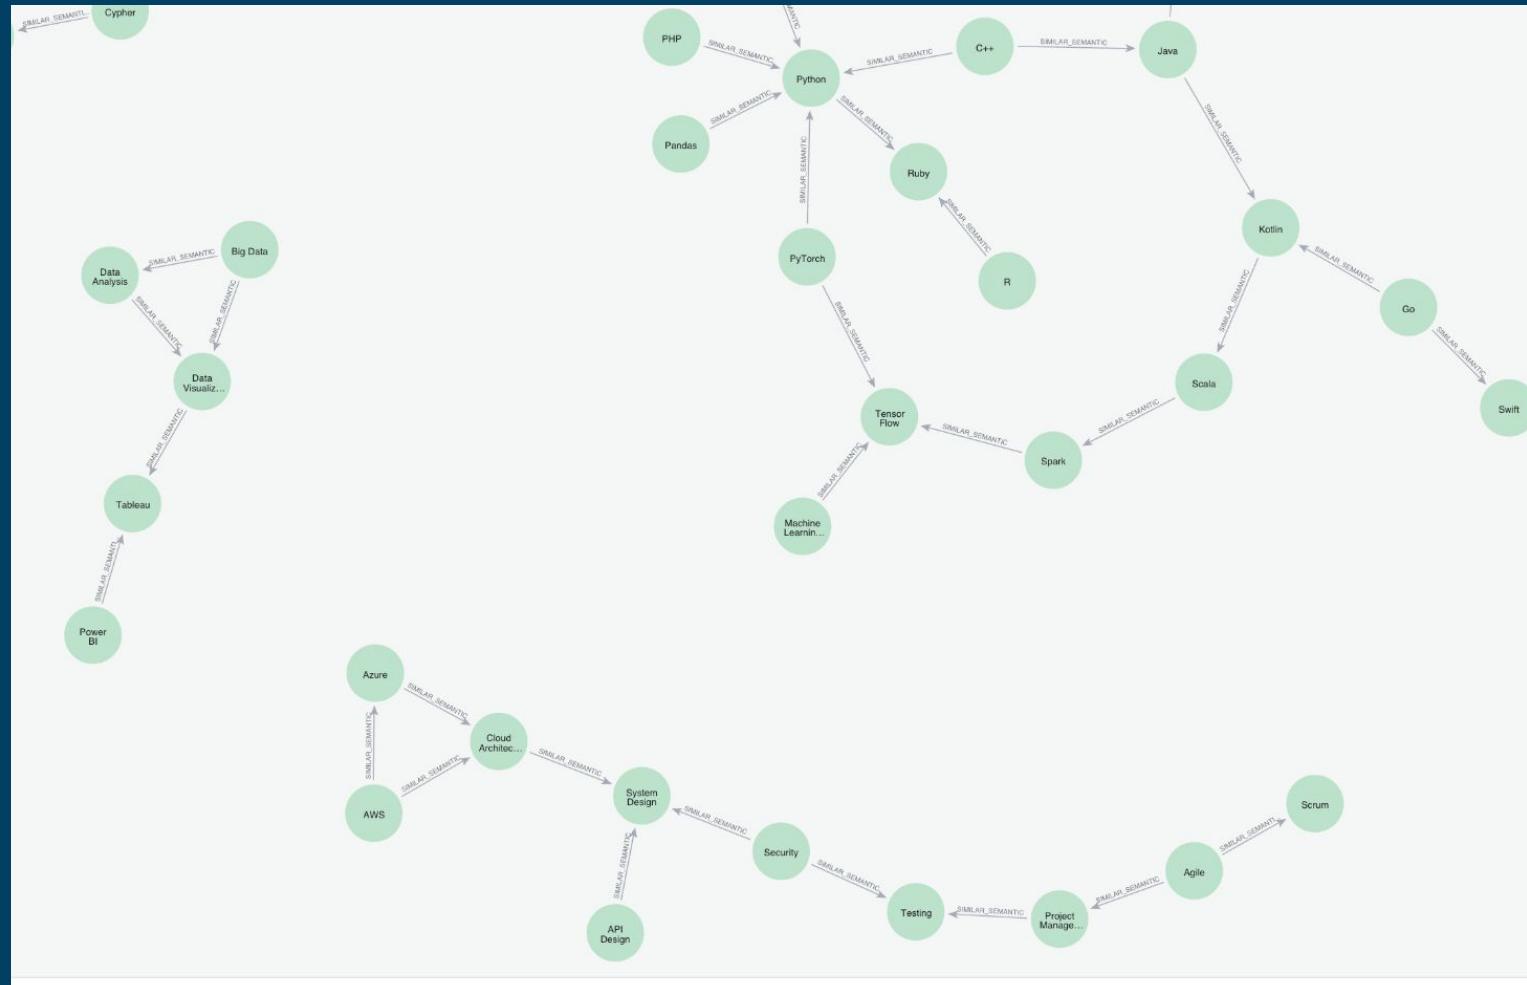
A Neo4j Knowledge Graph combined with LLM's obtain improvements:

- **Accuracy** - Obtain better answers compared to plain vector searches
- **Explainability**: Provide the user with more reasoning on how the results were obtained.
- **Acceleration**: Having all the capabilities in one platform increases understanding and decreases time to value.



Easier Development

Feedback from an AI Engineer



Here is the PR with the changes

I kinda replicated the same action-based cache already in place for Pinecone but thanks to the graph nature of Neo4J most of the operations yield better results:

- thanks to the [Neo4j graph data science](#) plugin we can store embeddings and calculate cosine similarity at the database level
- getting related actions is as simple as following the relationships between nodes
- **the cache can be visualised.** This is an extremely valuable debugging tool for us to understand if/when and how the cache might be broken/misbehaving (I actually already fixed a couple of bugs just thanks to this 🎉)

GraphRAG Resources



[Github Repository](#) with this workshop



Free online [graph academy](#) courses,
videos & webinars



Developer [guides and coded examples](#)

An aerial photograph of a school of dolphins swimming in the ocean. The water is a deep blue, and the dolphins are dark grey or black. They are swimming in several distinct groups, creating a sense of movement and social structure. The background shows the texture of the ocean surface.

Thank you!